INFO 324 Operating Systems II

Ahmad FAOUR

Virtual Memory (1)



When and why do we access a page table?

- On every instruction to translate virtual to physical addresses?

Page Talet

When and why do we access a page table?

- On every instruction to translate virtual to physical addresses?

In real machines NO!

In real machines it is only accessed

- On TLB miss faults to refill the TLB
- During process creation and destruction
- When a process allocates or frees memory?

Translation Lookaside Buffer

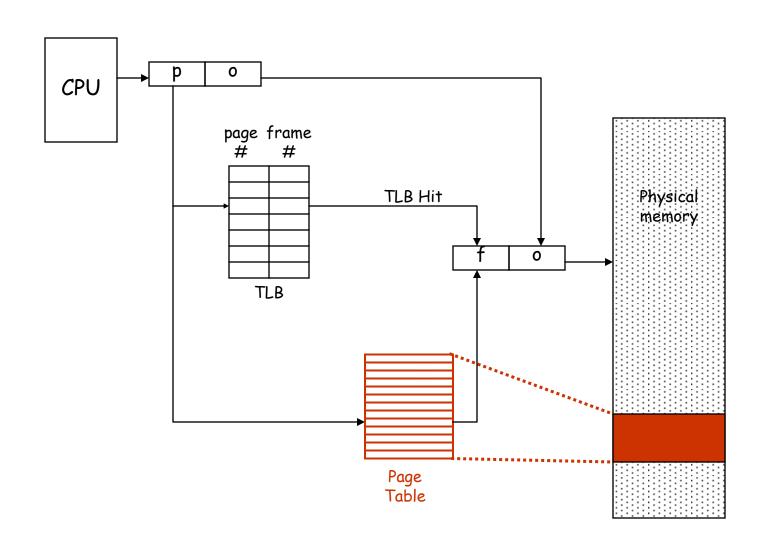
Problem: MMU can't keep up with the CPU if it goes to the page table on every memory access!

Cransation Leokaside Buffer

Solution:

- Cache the page table entries in a hardware cache
- Small number of entries (e.g., 64)
- Each entry contains page number and other stuff from page table entry
- Associatively indexed on page number
 ie. You can do a lookup in a single cycle

Eranslation Lookaside Buffer



Blandware Operation of TLB

K	e	У
		_

Page Number	Frame Number	Other	_			
23	37	unused	D	R	W	V
17	50	unused	D	R	M	V
92	24	unused	D	R	M	V
5	19	unused	D	R	M	V
12	6	unused	D	R	M	V

Hardware Operation of TLB



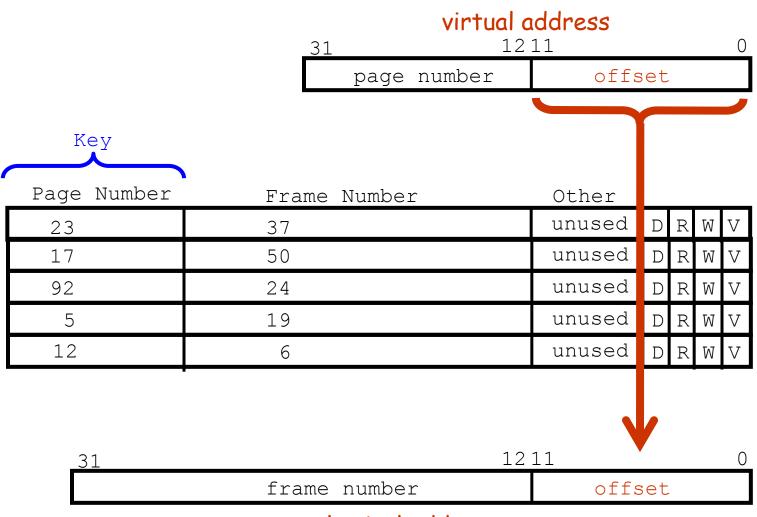
31	12 11			
	page	number	offset	

K	e	У
		_

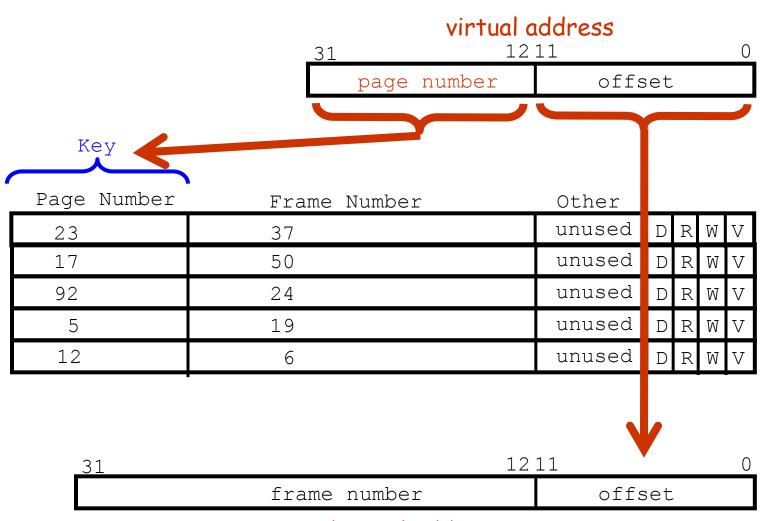
Page Number	Frame Number	Other				
23	37	unused	D	R	M	V
17	50	unused	D	R	W	V
92	24	unused	D	R	M	V
5	19	unused	D	R	W	V
12	6	unused	D	R	M	V

31	12 11			0
	frame number		offset	

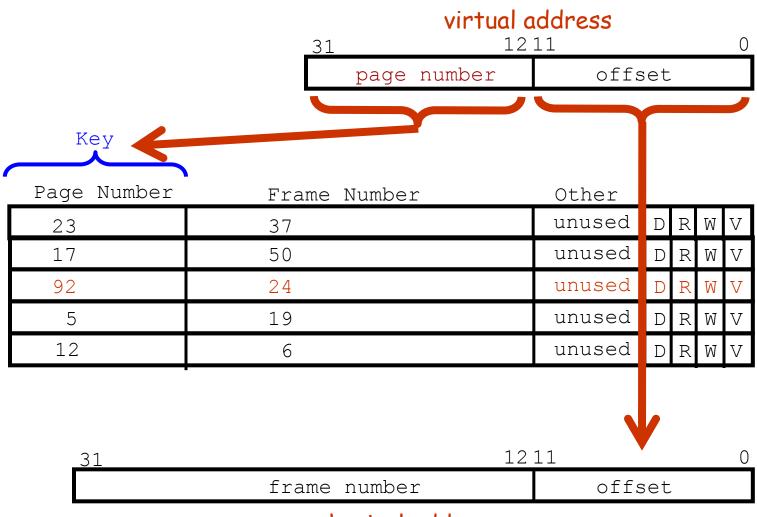
Elardware Operation of TLB



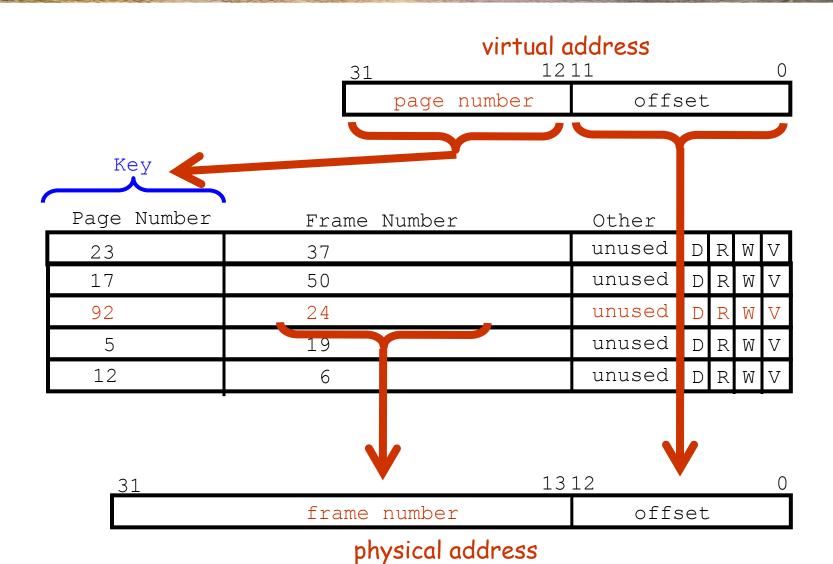
Hardware Operation of TLB



dardware Operation of TLB



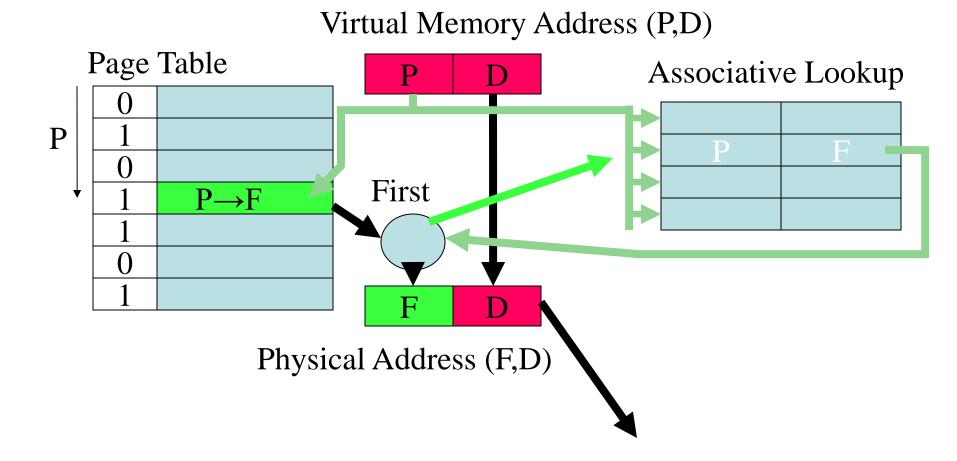
Elardware Operation of TLB



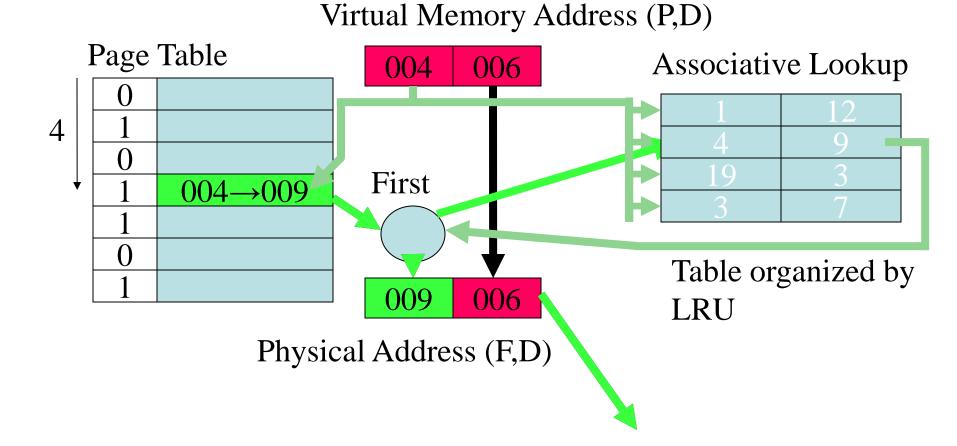
Beatchon

- If a virtual address is presented to MMU, the hardware checks TLB by comparing all entries simultaneously (in parallel).
- If match is valid, the page is taken from TLB without going through page table.
- If match is not valid
 - MMU detects miss and does an ordinary page table lookup.
 - It then evicts one page out of TLB and replaces it with the new entry, so that next time that page is found in TLB.

Page Mapping Hardware

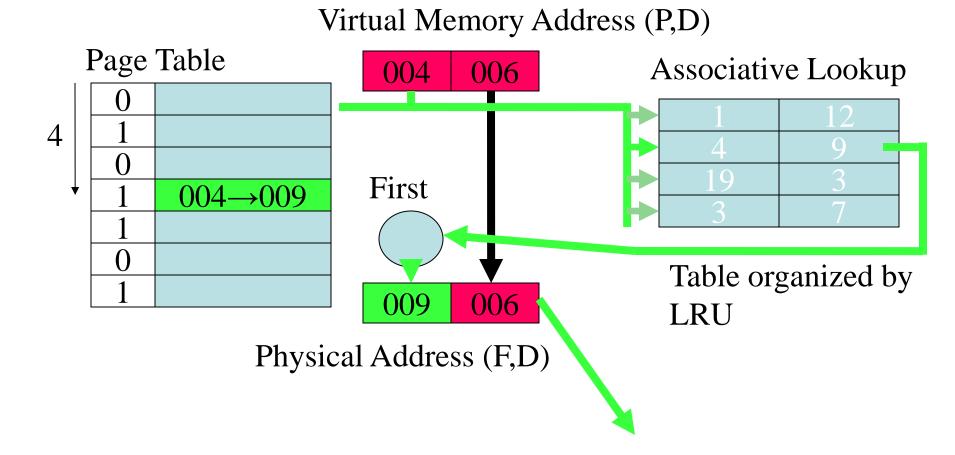


Page Mapping Example



INFO 324 - System Programming-2

Page Mapping Example: next reference



INFO 324 - System Programming-2

618 in a linery

Common (necessary) bits

Virtual page number: match with the virtual address

Physical page number: translated address

Valid

Access bits: kernel and user (nil, read, write)

Optional (useful) bits

Process tag

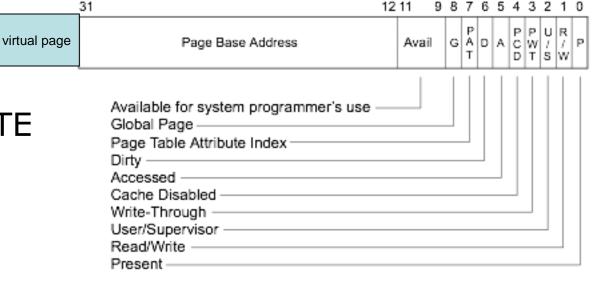
Reference

Modify

Cacheable

Includes parts of PTE

Example: x86



INFO 324 - System Programming-2

Page Talate Design

Page table size depends on

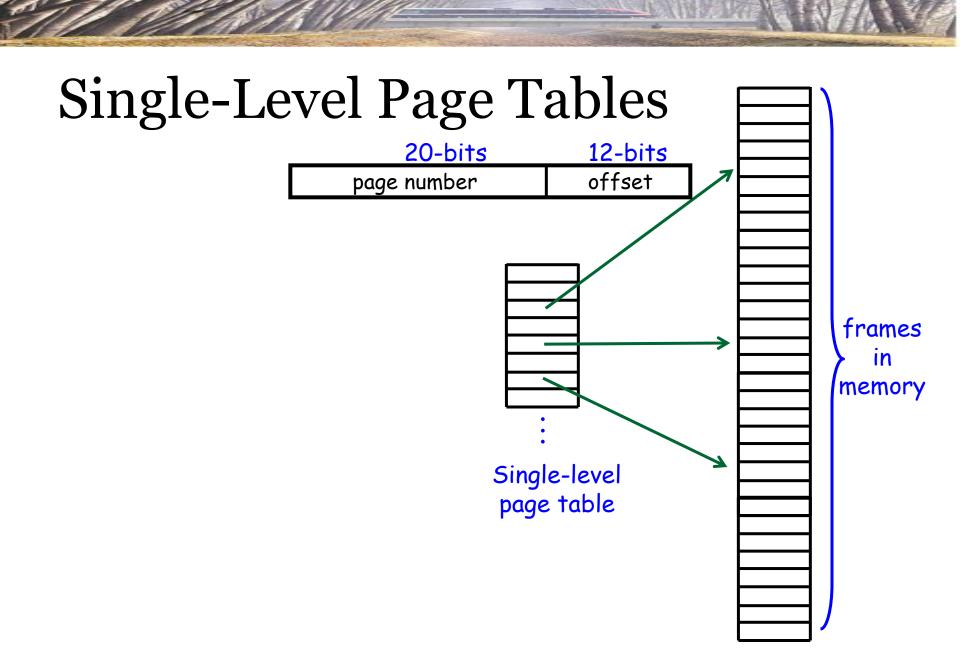
- Page size
- Virtual address length

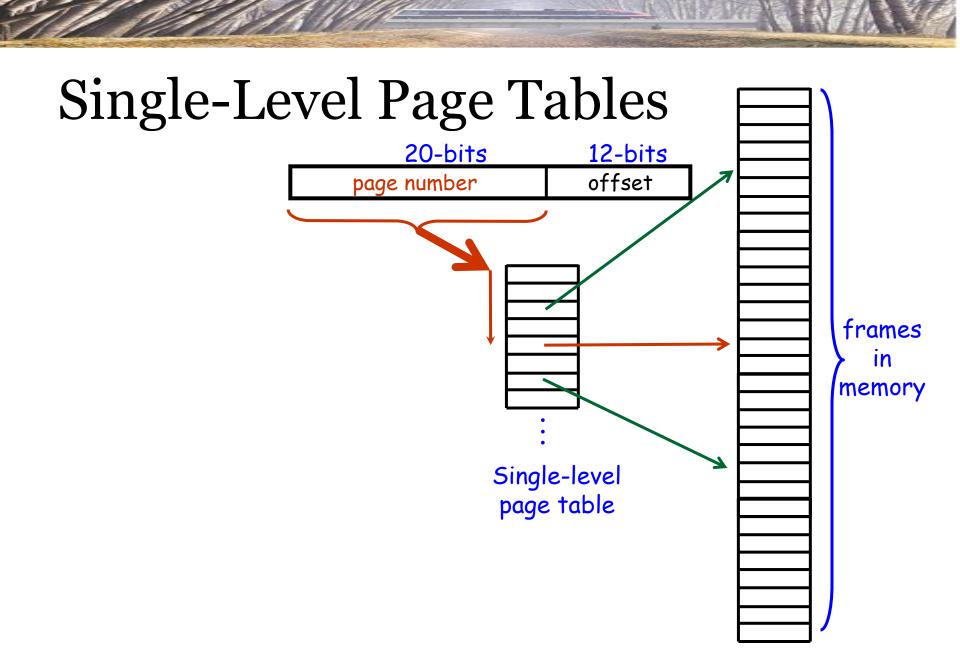
Memory used for page tables is overhead!

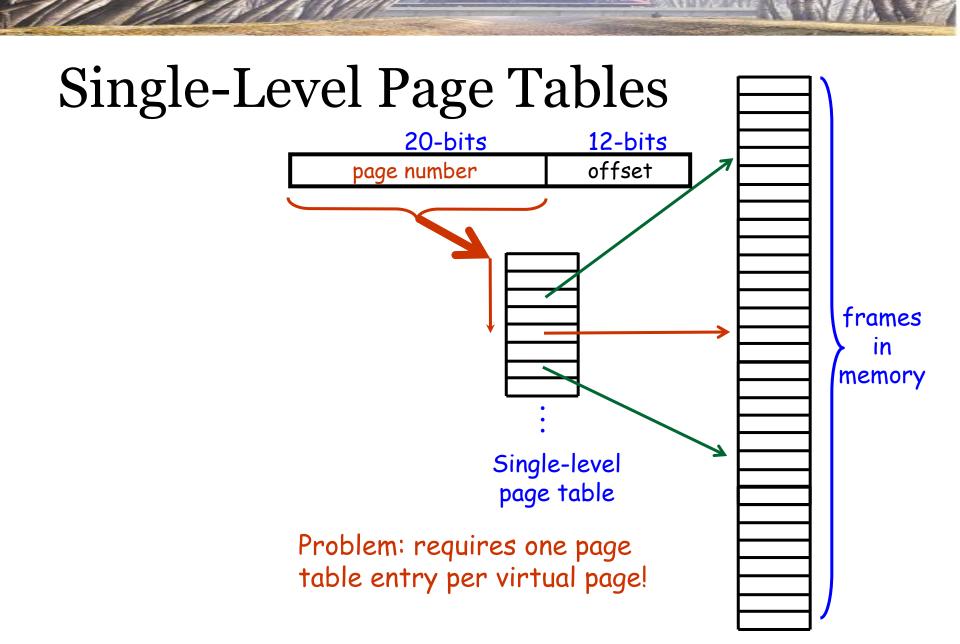
- How can we save space? ... and still find entries quickly?

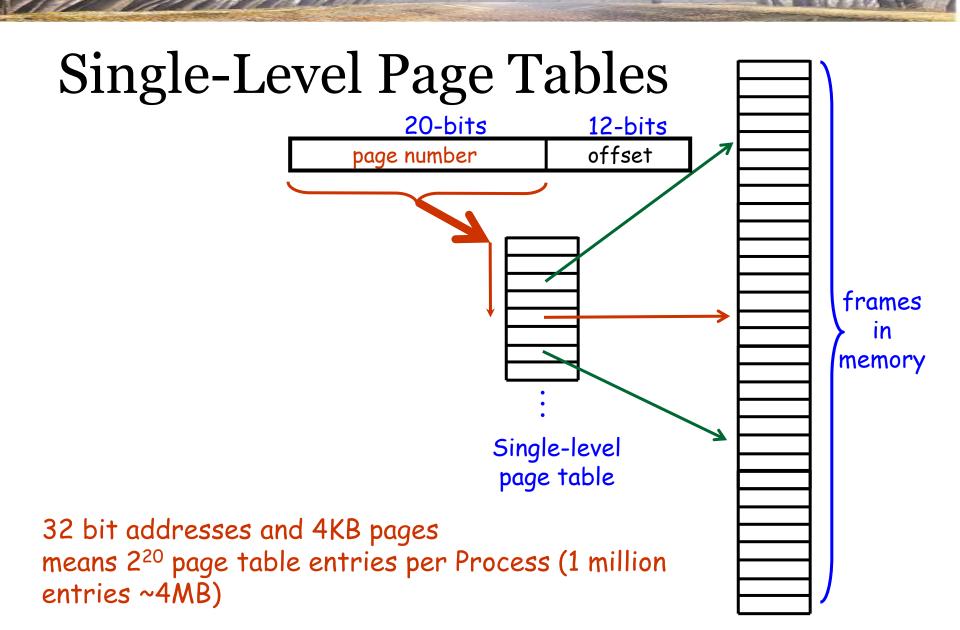
Three options

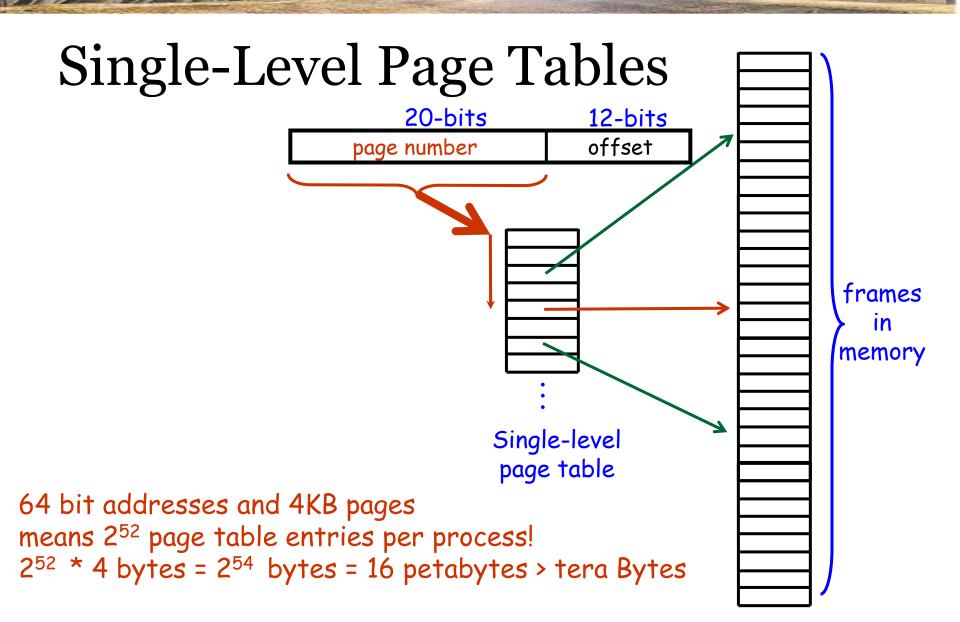
- Single-level page tables
- Multi-level page tables
- Inverted page tables





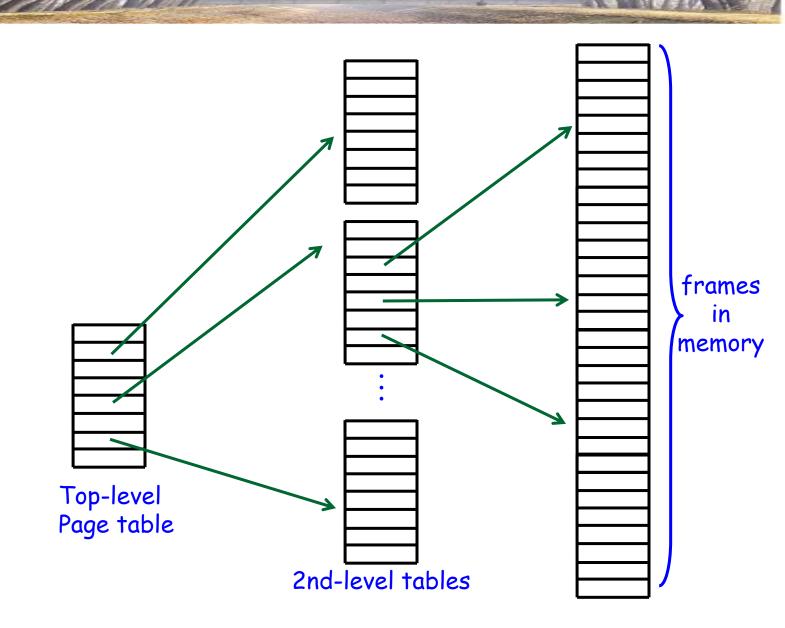






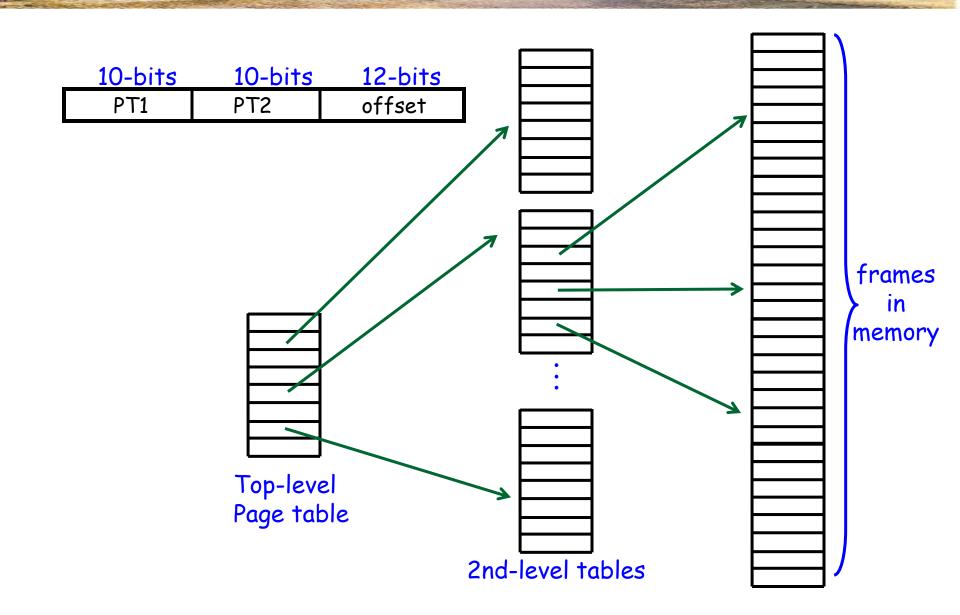
Maltileve Page Tables

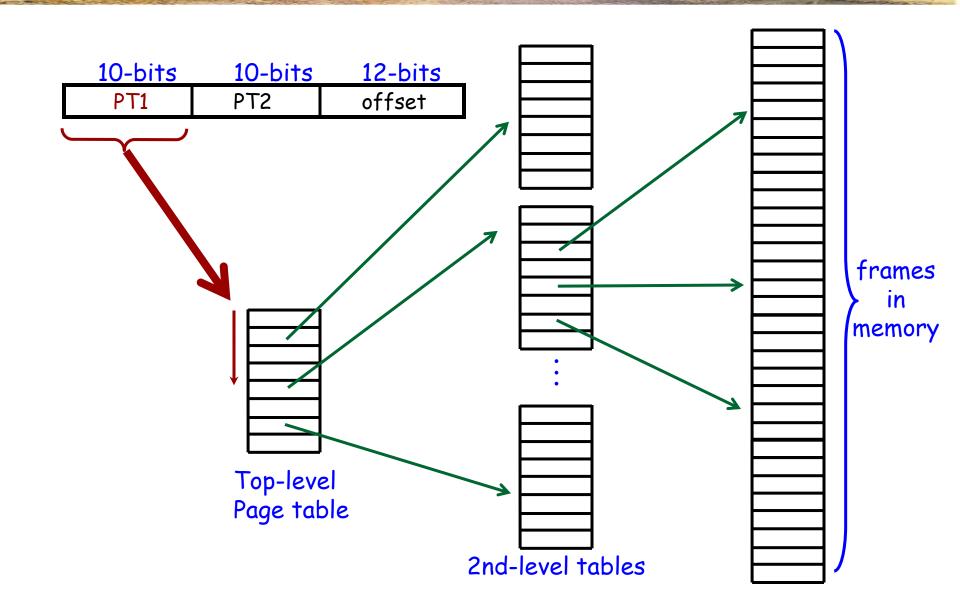
- Since the page table can be very large, one solution is to page the page table
- Divide the page number into
 - An index into a page table of second level page tables
 - A page within a second level page table
- Advantage
 - No need to keep all page tables in memory all the time
 - Only recently accessed memory's mapping need to be kept in memory, the rest can be fetched on demand

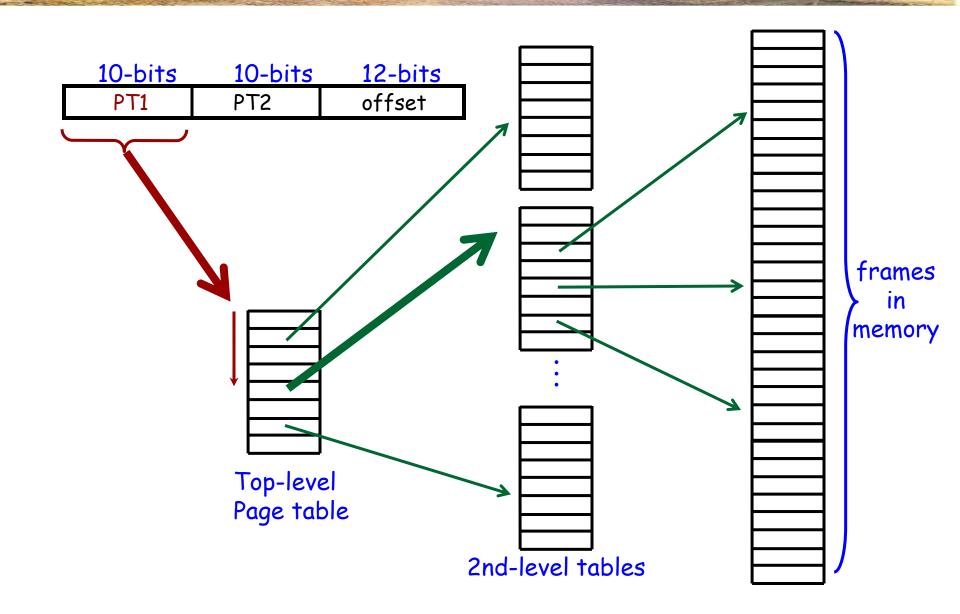


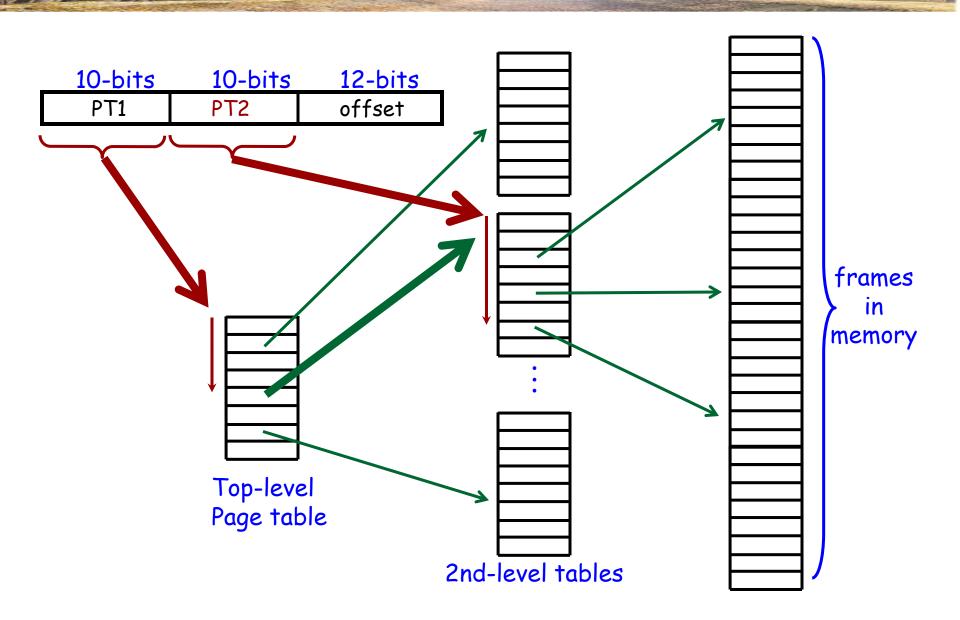
Multilevel page tables

- Multilevel page tables avoid keeping one huge page table in memory all the time.
- This works because most processes use only a few of its pages frequently and the rest, seldom if at all. Scheme: the page table itself is paged.
- o EX. Using 32 bit addressing:
 - o The top-level table contains 1,024 pages (indices).
 - The entry at each index contains the page frame number of a 2ndlevel page table.
 - This index (or page number) is found in the 10 highest (leftmost) bits in the virtual address generated by the CPU.
 - The next 10 bits in the address hold the index into the 2nd-level page table. This location holds the page frame number of the page itself.
 - The lowest 12 bits of the address is the offset, as usual.

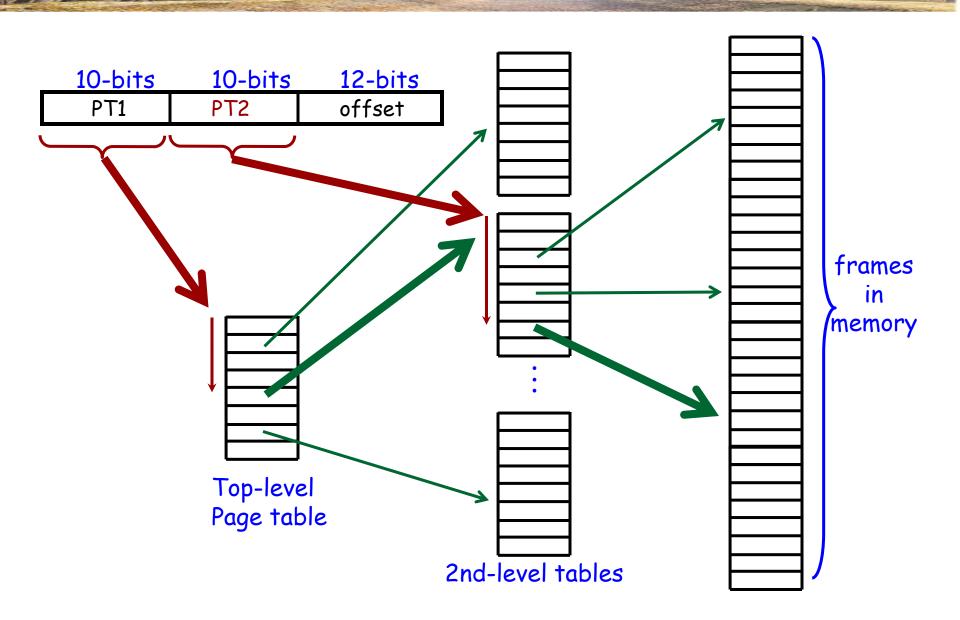


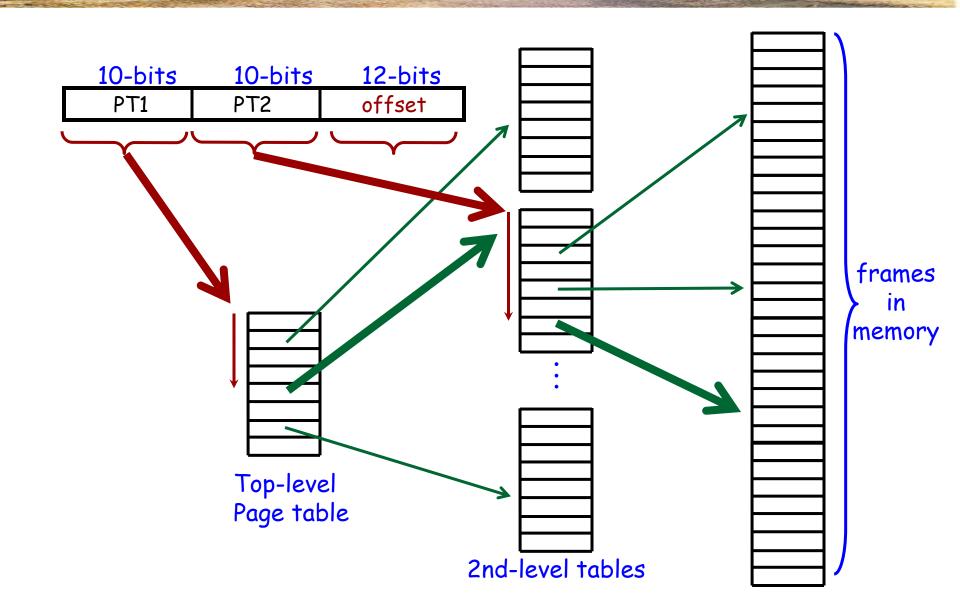






Maila-Level Page Tables





Ewo-taye Page Taletes (Example)

Ex. Given 32 bit virtual address 00403004 (hex) = 4,206,596 (dec)

converting to binary we have:

0000 0000 0100 0000 0011 0000 0000 0100

regrouping 10 highest bits, next 10 bits, remaining 12 bits:

0000 0000 01

00 0000 0011

0000 0000 0100

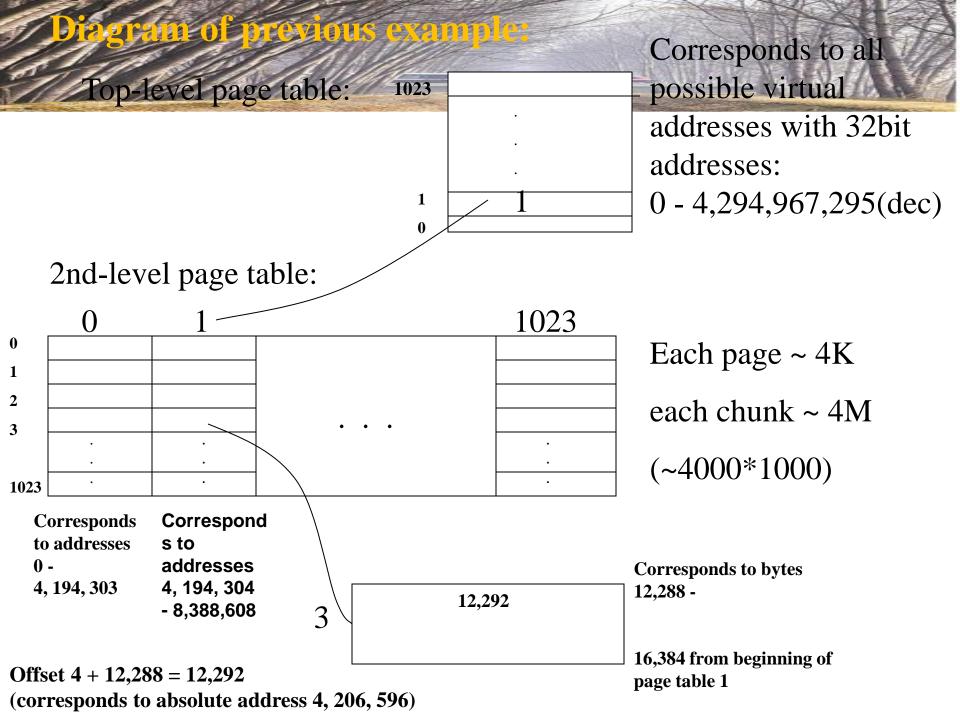
PT1 = 1

PT2 = 3

offset = 4

 $PT1 = 1 \Rightarrow$ go to index 1 in top-level page table. Entry here is the page frame number of the 2nd-level page table. (entry =1 in this ex.)

PT2 = 3 => go to index 3 of 2nd-level table 1. Entry here is the no. of the page frame that actually contains the address in physical memory. (entry=3 in this ex.) The address is found using the offset from the beginning of this page frame. (Remember each page frame corresponds to 4096 addresses of bytes of memory.)





Ok, but how exactly does this save space?

Ok, but how exactly does this save space?

Not all pages within a virtual address space are allocated

- Not only do they not have a page frame, but that range of

virtual addresses is not being used

- So no need to maintain complete information about it
- Some intermediate page tables are empty and not needed

We could also page the page table

- This saves space but slows access ... a lot!