

Part I

(29 pts)

A) Answer the following questions with justification :

- Explain the difference between segmentation and pagination, the relative interest of each other
- What is virtual memory? give the reasons why it is advantageous to have a virtual memory mechanism on a computer
- What is a fragmentation and how to remedy it in contiguous memory allocation

B) Consider the following C program:

1. Draw the graph generated by calling f(3).  
void f(int i){  
if(i<0) exit(0);  
if(i%2==0) fork();  
f(i-1);  
}

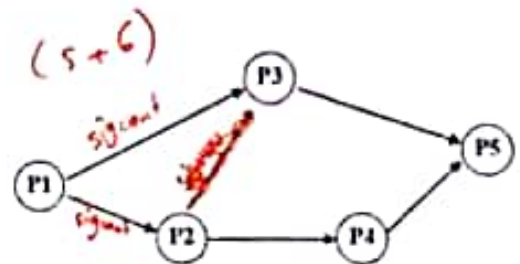
2. Draw the graph generated by the following program

void main(){  
for(int i = 0 ; i <= 3; i ++)  
f(i);  
}

3. Add (without any modification) to the function f the required statements such that no zombie processes are generated in part 2.

C) A parent process creates 5 child processes P1, P2, ..., P5. Then each child process displays its PID in respect to the order in the figure: For example the process P2 can't display its PID before P1. Also P5 can't before P3 and P4 and so on.

- Write the program using pipes of communication.
- Rewrite the same program using signals.



Part II

(18 pts)

A) Consider a paginated memory with size of 48 KB and pages of size 12 KB. The following references are requested in memory: A, B, C, D, E, B, E, F, D, A, B, C, G, F, C, B, A, B, C, F. How many page faults are generated using the following replacement algorithms?

- LRU (Least Recently Used)?
- Second chance algorithm?

B) We wish to allocate memory space for a process of 16KB.

- Simulate the functioning of First Fit and Best Fit algorithms on the following mapping.
- What gives the Next Fit algorithm if the former allocated block is indicated by the arrow X



C) We consider the following table of segments for a process P1

(6 pts)

Segment	Base	Limit
0	540	234
1	1254	128
2	54	328
3	2048	1024
4	976	200

- 1) Calculate the real addresses corresponding to the following virtual addresses (you may report addressing errors): (0:128), (1:100), (2:465), (3:888), (4:100), (4:344)
- 2) Is the virtual address (4,200) valid?

Note: the format of address is (segment#, offset)

### Part III (25 pts)

A) Suppose that a disk drive has 10,000 cylinders, numbered 0 to 9999. The driver is currently serving a request at cylinder 1400. The queue of pending requests is, in the order received:

100, 1200, 900, 8000, 8100, 100, 8200, 1000, 4200

- Starting from the current head position, what is the total distance (in cylinders) that the disk arm moves to satisfy all the pending requests for the following scheduling algorithms? (For the algorithms in which the head is in constant motion, indicate the direction in which you assume it is moving initially.)
- (a) FCFS (b) SSTF (c) Scan (d) C-look

B) We consider a file system that uses i-nodes like UNIX with few modifications as follows:

- 3 fields each of 8 bits containing information about the file
- 11 directs pointers to data blocks
- One pointer to simple indirection block where the last pointer of this block make another simple indirection

Given that each block has 1 Kb of size and occupies 2 bytes.

- a) What is the maximum size of a file in this system?
- b) Describe by figure the reading of the byte number 20992 of a file stored on disk.

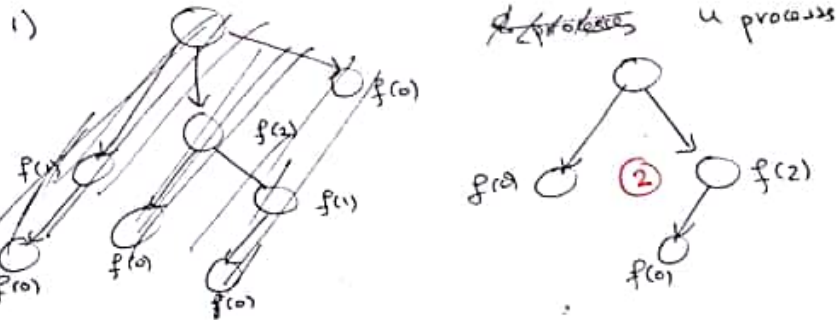
C) Refer to the function written in class (i.e., file\_open, ...):

- a. Describe (without writing code) the steps needed to open a file.
- b. How many I/O disk request is required to perform this task

Part I (27 pts:  $A \rightarrow 8 (1+2+3)$   
 $B \rightarrow 8 (2+5+1)$   
 $C \rightarrow 11 (5+6)$

B) void f(int i)  
 {  
   if (i < 0) exit(0);  
   if (i % 2 == 0) fork();  
   f(i-1);  
 }

$f(3) \rightarrow f(2) \xrightarrow{\text{fork}} f(1) \xrightarrow{\text{fork}} f(0) \xrightarrow{\text{fork}} f(-1) \times$



2) void main()  
 {  
   for (int i=0; i <= 3; i++)  
     f(i);  
 }

$f(0) \xrightarrow{\text{fork}} f(-1) \times$   
 $f(1) \rightarrow f(0)$   
 $f(2)$   
 $f(3)$

$\Rightarrow$  ~~fork~~ process

$i=0 \quad f(0) \xrightarrow{\text{fork}} f(-1) \times$   
 $i=1 \quad f(1) \rightarrow f(0) \xrightarrow{\text{fork}} f(-1) \times$   
 $i=2 \quad f(2) \xrightarrow{\text{fork}} f(1) \rightarrow f(0) \xrightarrow{\text{fork}} f(-1) \times$   
 $i=3 \quad f(3) \rightarrow f(2) \xrightarrow{\text{fork}} f(1) \rightarrow f(0) \xrightarrow{\text{fork}} f(-1) \times$   
 we have 6 calls to fork, so  $2^6$  processes = **64** (5)

2



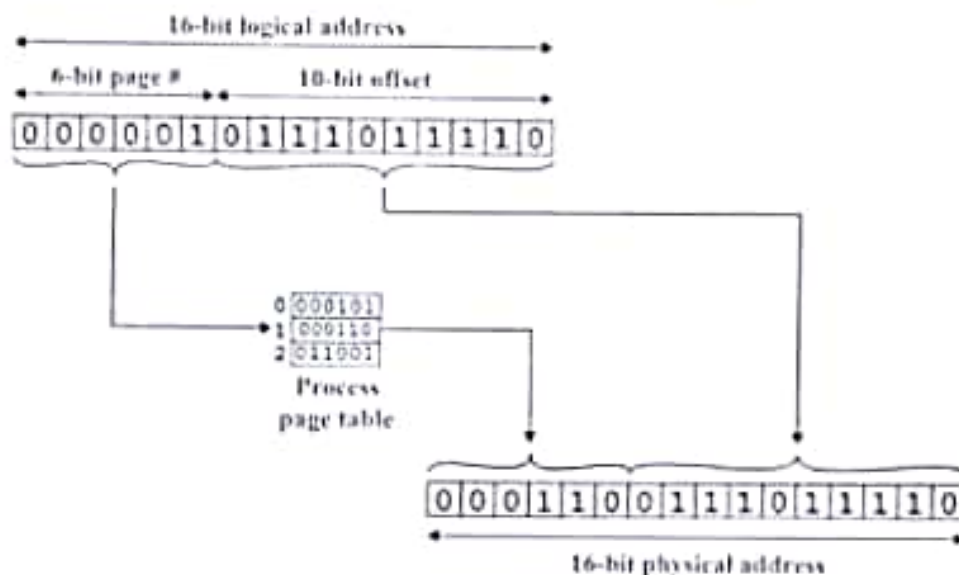
## Part 1

a) Difference between Segmentation and Pagination: (3 pts)

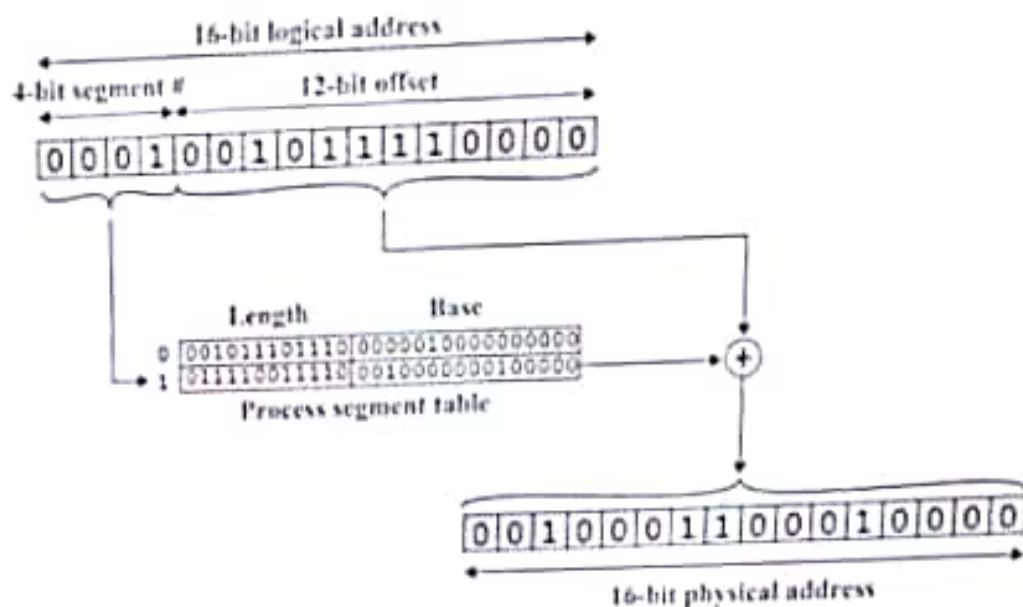
- Paging is used to get a large linear address space without having to buy more physical memory. Segmentation allows programs and data to be broken up into logically independent address spaces and to aid sharing and protection.
- Paging does not distinguish and protect procedures and data separately.
- Segmentation distinguishes and separately protects procedures and data.
- Unlike segmentation, Paging does not facilitate sharing of procedures.
- Paging is transparent to programmers (system handles it automatically)
- Segmentation requires programmer to be aware of memory limits as programmer tries to allocate memory to functions and variables or tries to access read only memory violation, which results in segmentation fault.
- Mapping from logical to physical address is different for paging and segmentation.  
Here's an illustration based on 16 bit-address space:

**For paging:**

The 6-bit page value is used to select a proper entry in process page table. the 6-bit process entry occupying the six most significant bit and the 10-bit offset occupying the 10 least significant bit forms a 16-bit physical address.

**For segmentation:**

The 4-bit segment of a logical address selects the proper entry in the process segment table. The base value is added to the 12 bit offset value to get the 16 bit physical address.



b) **What is virtual memory?** (2 pts)

In computing, virtual memory is a memory management technique that is implemented using both hardware and software. It maps memory addresses used by a program, called **virtual addresses**, into physical addresses in computer memory. Main storage as seen by a process or task appears as a **contiguous address space or collection of contiguous segments**. The operating system manages virtual address spaces and the assignment of real memory to virtual memory. Address translation hardware in the CPU, often referred to as a memory management unit or MMU, automatically translates virtual addresses to physical addresses. Software within the operating system may extend these capabilities to provide a virtual address space that can exceed the capacity of real memory and thus reference more memory than is physically present in the computer.

The primary benefits of virtual memory include freeing applications from having to manage a shared memory space, increased security due to memory isolation, and being able to conceptually use more memory than might be physically available, using the technique of paging.

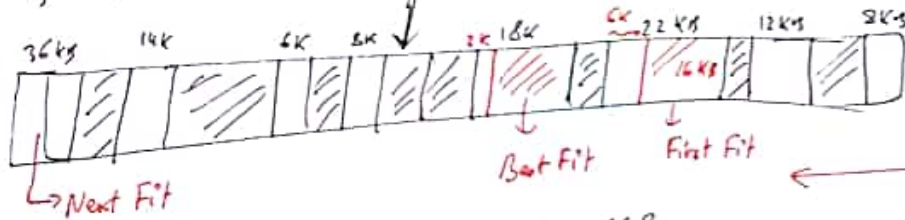
c) **What is a fragmentation?** (3 pts)

As processes are loaded and removed from memory, the free memory space is broken into little pieces. It happens after sometimes that processes cannot be allocated to memory blocks considering their small size and memory blocks remains unused. This problem is known as Fragmentation.

Fragmentation is of two types

B) process = 16 KB

1) First Fit 2. best Fit



- C) 1)  $(0:128) \rightarrow 540 + 128 = 668$   
 $(1:100) \rightarrow 1254 + 100 = 1354$   
 $(2:465) \rightarrow \text{error}$   
 $(3:888) \rightarrow 2048 + 888 = 2936$   
 $(4:100) \rightarrow 976 + 100 = 1076$   
 $(4:344) \rightarrow \text{error}$   
 2)  $(4, 200)$  NO

Part III Disk: 10000 cylinders

A) - currently at cylinder 1400  
 queue: 100, 1200, 900, 8000, 2100, 100, 8200, 1000,  
 4200

- a) FCFS  $\rightarrow 36400$   
 b) SSTF  $\rightarrow 9400$   
 c) Scan  $\rightarrow 9600$  or  $18498$   
 d) C-look  $\rightarrow 13400$  or  $16000$

5300  
 x 9100  
 5300

```

3) void f(int i)
{
    if (i < 0) exit(0);
    if (i % 2 == 0)
    {
        fork();
        wait(NULL);
    }
    f(i-1);
}

```

c) #include <stdio.h>  
#include <stdlib.h>  
#include <unistd.h>

```

void main() {
    int p1[2], p2[2], p3[2],
    p4[2], p5[2];

```

```

    int x;
    pipe(p1); pipe(p2); .....
    close(p1[0]);
    for (int i=1; i<5; i++)
    {
        if (!fork()) break;

```

```

        if (i==1) { // P1
            printf("Process P1 with pid %d\n", getpid());
            write(p1[1], &x, sizeof(int));
            write(p2[1], &x, sizeof(int));
        }
        else if (i==2) // P2
        {
            read(p1[0], &x, sizeof(int));
            printf("Process P2 ---");
            write(p2[1], &x, sizeof(int));
        }
        else if (i==3) // P3

```

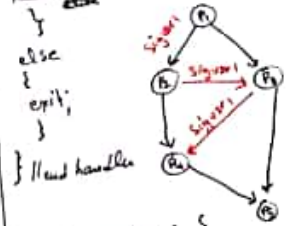
5 pts

2) signal

```

#include <stdio.h>
#include <stdlib.h>
#include <sys/wait.h>
int pid, int fd[2]; int next;
static counter = 0;
void handler (int sig)
{
    ctr++;
    printf("I'm the process number %d, with pid %d\n", ctr, getpid());
    if (read(fd[0], &next, sizeof(int)))
    {
        kill(next, sigusr1); exit();
    }
    else
    {
        exit();
    }
} // end handler

```



5 pts

```

void main() {
    int pipe(fd);
    signal(sigusr1, handler);
    for (int i=1; i<5; i++)
    {
        if (!pid = fork())
        {
            close(fd[1]);
            pause();
            write(fd[1], &pid, sizeof(int));
        } // end for
        printf("I'm the process P1 with pid = %d\n", getpid());
        close(fd[1]);
        while (wait(NULL));
    } // end main
    read(fd[0], &next, sizeof(int));
    kill(next, sigusr1);
    counter++;
}

```

(3)

(3)



## Second chance Alg

Page	Point	6 pages from			
A	yes	A			
B	yes	B	A		
C	yes	C	B	A	
D	yes	D	C	B	A
E	yes	E	D	C	B
B	No	E	D	C	B*
E	No	E*	D	C	B*
F	yes	F	E*	D	B
D	No	F	E*	D*	B
A	yes	A	F	E*	D*
B	yes	B	A	E	D
C	yes	C	B	A	E
G	yes	G	C	B	A
F	yes	F	E	C	B
C	No	F	G	C*	B
B	No	F	G	C*	B*
A	yes	A	F	C	B
B	No	A	F	C	B*
C	No	A	F	C*	B*
F	No	A	F*	C*	B*

12

12 page faults



## Part II

A) Memory: size 48 KB }  $\Rightarrow$  4 Frames  
 - page 12 KB

a) LRU (5) (3+2)

Page ref	Faults				
A	yes	A			
B	yes	A	B		
C	yes	A	B	C	D
D	yes	A	B	C	D
E	yes	E	B	C	D
B	No	E	B	C	D
E	No	E	B	C	D
F	yes	E	B	F	D
D	No	E	B	F	D
A	yes	E	A	F	D
B	yes	B	A	F	D
C	yes	B	A	C	G
G	yes	B	A	C	G
F	yes	B	F	C	G
C	No	B	F	C	G
B	No	B	F	C	A
A	yes	B	F	C	A
B	No	B	F	C	A
C	No	B	F	C	A
F	No	B	F	C	A

12 page faults

b) Second chance (5) (3+2)

Page ref	4 Page Frames				
	Faults	Page contents			
A	Yes	A*			
B	yes	B*	A*		
C	yes	C*	B*	A*	
D	yes	D*	C*	B*	A*
E	yes	E*	D	C	B
B	No	E*	D	C	B*
E	No	E*	D	C	B
F	yes	F*	D	C	B
D	No	D*	E*	F*	B
A	yes	A*	D*	E*	F*
B	yes	B*	A	F	E
C	yes	C*	B*	A	F
G	yes	G*	C*	B*	A
F	yes	F*	G*	C*	B*
C	No	F*	G*	C*	B*
B	No	F*	G*	C*	B*
A	yes	A*	F	E	B
B	yes	B*	A*	F	C
C	yes	C*	B*	F	C
F	No	F*	C*	B*	A*

✓ 12 page faults  
 12  $\rightarrow$  see next page  
 (3)

c) file-open ( $10 = 6 + 4$ )

a) steps

- 1- for opening the file we need to find its associated inodo #
- 2- search the global folder for the (external name)
  - 2.1 research loop into the entry 0 of (fdesc)
  - 2.2 load the data blocks for the global folder into memory (buffer)
  - 2.3 each 16 bytes corresponds to (name, inodo)
- 3- when the file is found, get the inodo #
- 4- search for a free entry in the table of descriptor (fdesc)
- 5- load in memory the block of inodes that contains the inodo #
- 6- initialize the fdesc entries

b) I/O disks?

- \* file-set-position(0,0)  $\rightarrow$  1 disk read (I/O)
- \* in worst case we need to read all blocks of the global directory (entry 0): Suppose  $x$  I/O
- \* disk-read (binode,  $2 + (\text{fold.inode\_nb}/16)$ );  $\rightarrow 2$  I/O

So we need  $\boxed{x+2 \text{ I/O}}$

**Part 1: Process Management (30 points):**

- A) Answer the following questions with justification :
- Explain in detail how the use of two-level or more page tables will reduce the space required for the one-level page table. (2)
  - Explain the difference between logical and physical addresses (3)
- B) After explaining what it means:
- A process is in the zombie state. (2)
  - A process is orphaned. (5)
- Write a program that creates a process that fall within the zombie state before becoming orphan.
- C) Write a C program where a parent creates N child processes. Then, the processes behave in the following manner:  
Every T second, the father write in a pipe a character 'G' and sends to all "alive" child the "SIGUSR1" signal for unlock them. The child awakened by the signal, start the race to read the character already written in the tube by the father. The winning child process must terminate after displaying a number indicating the termination order.  
P.S: Remember that the father sent the signal "SIGUSR1" only for living child (who are not already terminated.) (15)

**Part 2: Memory Management (15 points)**

Consider a system with paginated main memory. The memory is composed of 4 frames (frames each has a size of 64 bytes. At a given moment, the memory is empty, then two processes P1 (4 pages) and P2 (2 pages) are launched in the system. The processor sends requests submissions in the following order in the format [hexadecimal logical address, Process]:

[3F, P1]  
[4A, P1]  
[1D, P1]  
[00, P2]  
[CA, P1]  
[87, P1]  
[39, P2]  
[2B, P1]  
[00, P1]  
[11, P2]

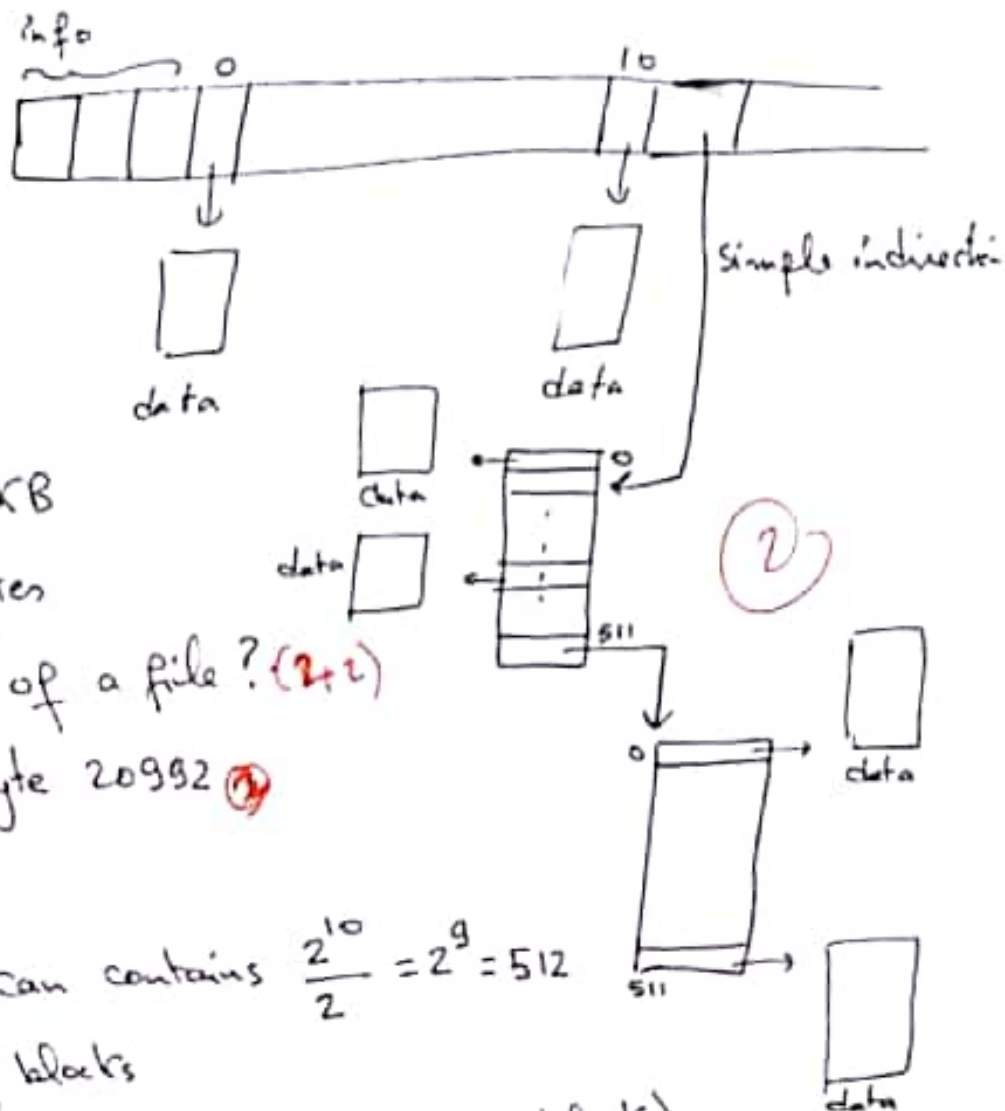
- What is in bytes the size of the main memory? (3)
- What is the size of the virtual address space? (3)
- Indicate by figures the evolution of the memory and the number of page faults using the following page replacement algorithms :
  - FIFO (3)
  - LRU (2)
  - Second chance. (1)



# Part III

B) Inode :

(7) (4+1)



Block size = 1 KB

Block # = 2 bytes

a) Maximum size of a file? (2+2)

b) reading of byte 20992 (2)

Answers :

a) each block can contains  $\frac{2^{10}}{2} = 2^9 = 512$  number of blocks

Max size : 11 direct block + 512 (data block)  
+ 512 (data block)

$$= 11 + 511 + 512 = 1034 \text{ KB}$$

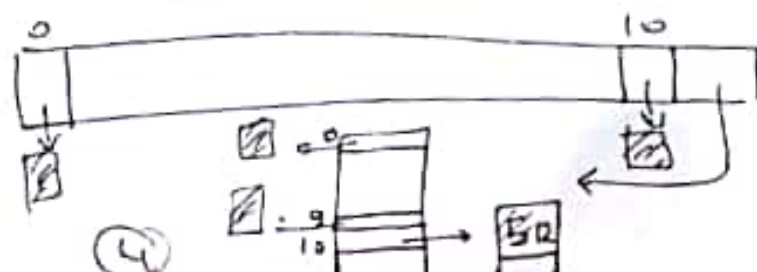
b)  $20992 - (11 \times 1024) = 20992 - 11264 = 9728$

$$9728 = (9 \times 1024) + 512$$

So the File covers the 11 direct data block  
- the first 0  $\rightarrow$  9 (10) from the simple indirect  
- In the block 10 (simple indirection), we have

offset 512

(3)



(4)

### Part 3: File System (25 points)

- A) Consider a file system with linked list and indexed allocation strategy as illustrated in the figure. The size of block is 8 KB.

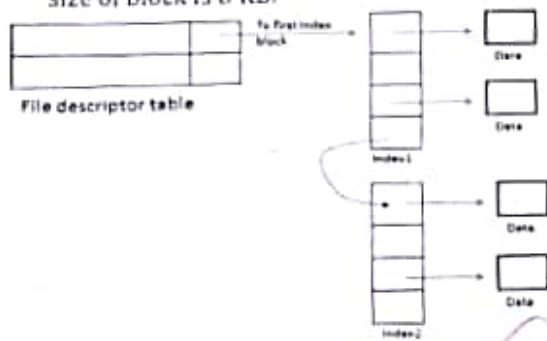


Figure: File system illustration

The disk blocks are numerated from 0 to  $n-1$ . Each opened file has an entry in the file descriptor table which is loaded in memory. This entry contains all attributes of the file and a pointer to an index block. The index block contains  $k$  pointers to other blocks. These pointers except the last points to data block. The last pointer points to another index block and so on as shown in figure. The pointer to block occupies 32 bits.

**Question:** What is the maximum size (in number of blocks) of a file in this system? Indicate the number of data blocks and index blocks

- B) Given a FS where the topo table contains 10 entries, each corresponding to a single level indexing (each block points to a map). Given that each block occupies 2 kilobytes, the number of a block occupies 4 bytes and each block contains 16 inode:  
Write the function ***delete\_block(int lbd, int f)*** that delete the physical data block corresponding to the logical data block lbd of the file with descriptor f.  
Note that if this block is not at the end of file, you must shift all the other data blocks that follow.

**P.S:** the use of the functions seen in the course is permitted if needed.

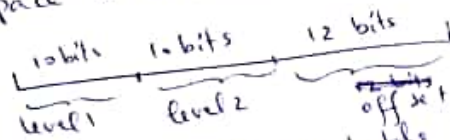
Part 1: Process managementa) a) using 2 or more level page table

Since each page in the virtual address space needs an entry in the page table, we may have a huge number of entries in the page table when actually a few of them are needed by the process. For example in a virtual address space of 32 bits with page of size 4KB =  $2^{12}$  we need  $2^{20}$  entries. Each entry occupies 4 bytes

$\Rightarrow$  4KB the page table are reserved for a process which it actually needs a few entries. For a 64 bits virtual address space we need  $2^{52} \times 4$  bytes for the page table for each process

$\rightarrow$  using multi-level page table can reduce the overhead as follows:

For example in two-level page table the virtual address space is divided into 3 parts: (32 bits)



there are two page table indexes:

- the first page index is used to index the outer page table
- each entry in the outer page table can possibly point to a second level page table
- the present bit in the outer page table is used to indicate if there is a second level page table for this entry

①

②



## Part 2: Memory management (15 pts)

$P_1 \rightarrow 4$  pages

$P_2 \rightarrow 2$  pages

1) size of entry:  
 $64 \times 4 = 256 \text{ Bytes} = 2^8 \text{ Bytes}$  (3)

2) size of the virtual address space

~~Here~~ Page size =  $64 \text{ Bytes} = 2^6 \text{ Bytes}$

$\Rightarrow 6$  bits is needed for the offset

- The logical address is in hexadecimal (2 digits)  
 - each is encoded on 4 bits  $\Rightarrow$  the logical address is encoded on 8 bits  $\Rightarrow$   $\frac{11}{2} \times 6$

$\Rightarrow 2^2$  pages  $\Rightarrow 4 \times 64 = 256 \text{ Bytes} = 2^8$  (3)

Hex	Binary	Page #	Offset	Physical Address
3F	0011 1111	63	0	$P_1$
4A	0100 1010	74	1	$P_1$
1D	0001 1101	29	0	$P_1$
00	0000 0000	0	0	$P_2$
CA	1100 1010	202	3	$P_1$
27	1000 0111	135	2	$P_1$
39	0011 1001	57	0	$P_2$
2B	0010 1011	43	0	$P_1$
00	0000 0000	00	0	$P_1$
11	0001 0001	17	0	$P_2$

$\Rightarrow$  The  $W = \{0, 1, 0, 0, 3, 2, 0, 0, 0, 0\}$

$W = \{(0; P_1), (1; P_1), (0; P_1), (0; P_2), (3; P_1), (2; P_1), (0; P_2), (0; P_1), (0; P_2)\}$

a) FIFO

	0	1	2	3	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
1		1	1	1	1	1	1	1	1
2				3	3	3	3	3	3
3					2	2	2	2	2

(3)

### a) FIFO

Page refs	Fault?	4 page frames			
		Page contents			
(0, P <sub>1</sub> )	✓	(0, P <sub>1</sub> )			
(1, P <sub>1</sub> )	✓	(1, P <sub>1</sub> )	(0, P <sub>1</sub> )		
(0, P <sub>1</sub> )	X	(1, P <sub>1</sub> )	(0, P <sub>1</sub> )		
(0, P <sub>2</sub> )	✓	(0, P <sub>2</sub> )	(0, P <sub>1</sub> )	(0, P <sub>1</sub> )	
(3, P <sub>1</sub> )	✓	(3, P <sub>1</sub> )	(1, P <sub>1</sub> )	(0, P <sub>1</sub> )	
(2, P <sub>1</sub> )	✓	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>1</sub> )	(1, P <sub>1</sub> )
(0, P <sub>2</sub> )	X	(3, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>2</sub> )	(1, P <sub>1</sub> )
(0, P <sub>1</sub> )	✓	(0, P <sub>1</sub> )	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>1</sub> )
(0, P <sub>1</sub> )	X	-	-	-	-
(0, P <sub>2</sub> )	X	-	-	-	-

6 page faults (3)

### b) LRU

Page refs	Fault?	4 Page frames			
		Page contents			
(0, P <sub>1</sub> )	✓	(0, P <sub>1</sub> )			
(1, P <sub>1</sub> )	✓	(1, P <sub>1</sub> )	(0, P <sub>1</sub> )		
(0, P <sub>1</sub> )	X	(0, P <sub>1</sub> )	(1, P <sub>1</sub> )		
(0, P <sub>2</sub> )	✓	(0, P <sub>2</sub> )	(0, P <sub>1</sub> )	(1, P <sub>1</sub> )	
(3, P <sub>1</sub> )	✓	(3, P <sub>1</sub> )	(0, P <sub>2</sub> )	(0, P <sub>1</sub> )	(1, P <sub>1</sub> )
(2, P <sub>1</sub> )	✓	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>2</sub> )	(0, P <sub>1</sub> )
(0, P <sub>2</sub> )	X	(0, P <sub>2</sub> )	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>1</sub> )
(0, P <sub>1</sub> )	X	(0, P <sub>1</sub> )	(0, P <sub>2</sub> )	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )
(0, P <sub>1</sub> )	X	-	-	-	-
(0, P <sub>2</sub> )	X	(0, P <sub>2</sub> )	(0, P <sub>1</sub> )	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )

5 page faults (3)

### c) Second chance

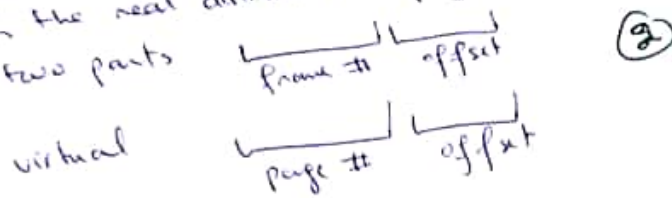
Page ref	4 Page Frames					out
	Fault?	Page contents				
(0, P <sub>1</sub> )	✓	(0, P <sub>1</sub> )				
(1, P <sub>1</sub> )	✓	(1, P <sub>1</sub> )	(0, P <sub>1</sub> )			
(0, P <sub>1</sub> )	X	(1, P <sub>1</sub> )	(0, P <sub>1</sub> ) <sup>*</sup>			
(0, P <sub>2</sub> )	✓	(0, P <sub>2</sub> )	(1, P <sub>1</sub> )	(0, P <sub>1</sub> ) <sup>*</sup>		
(3, P <sub>1</sub> )	✓	(3, P <sub>1</sub> )	(0, P <sub>2</sub> )	(1, P <sub>1</sub> )	(0, P <sub>1</sub> ) <sup>*</sup>	
(2, P <sub>1</sub> )	✓	(2, P <sub>1</sub> )	<del>(3, P<sub>1</sub>)</del>	(0, P <sub>2</sub> )	(0, P <sub>1</sub> )	(1, P <sub>1</sub> )
(0, P <sub>2</sub> )	X	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>2</sub> ) <sup>*</sup>	(0, P <sub>1</sub> )	
(0, P <sub>1</sub> )	X	(2, P <sub>1</sub> )	(3, P <sub>1</sub> )	(0, P <sub>2</sub> ) <sup>*</sup>	(0, P <sub>1</sub> ) <sup>*</sup>	
(0, P <sub>1</sub> )	X	—	—	—	—	
(0, P <sub>2</sub> )	X	—	—	—	—	

5 page faults

(3)

## b) logical & physical address

The logical address is used by each process in ~~each~~ its virtual address space starting from 0 to max-1 bytes depending on the architecture and it is private for each process. whereas the physical address ~~space~~ is the real address on physical memory. It has two parts



## B) a) zombie state

a zombie process is a process that has completed execution (via the exit system call) but still has an entry in the process table waiting his parent to get its status by the (wait call)

## b) orphan state

~~process~~ An orphan process is a process that is still executing, but whose parent has died

```
→ #include <unistd.h>
   #include <stdio.h>
```

```
int main()
```

```
{ int pid; int i=0;
```

```
  pid = fork();
```

```
  if (!pid)
```

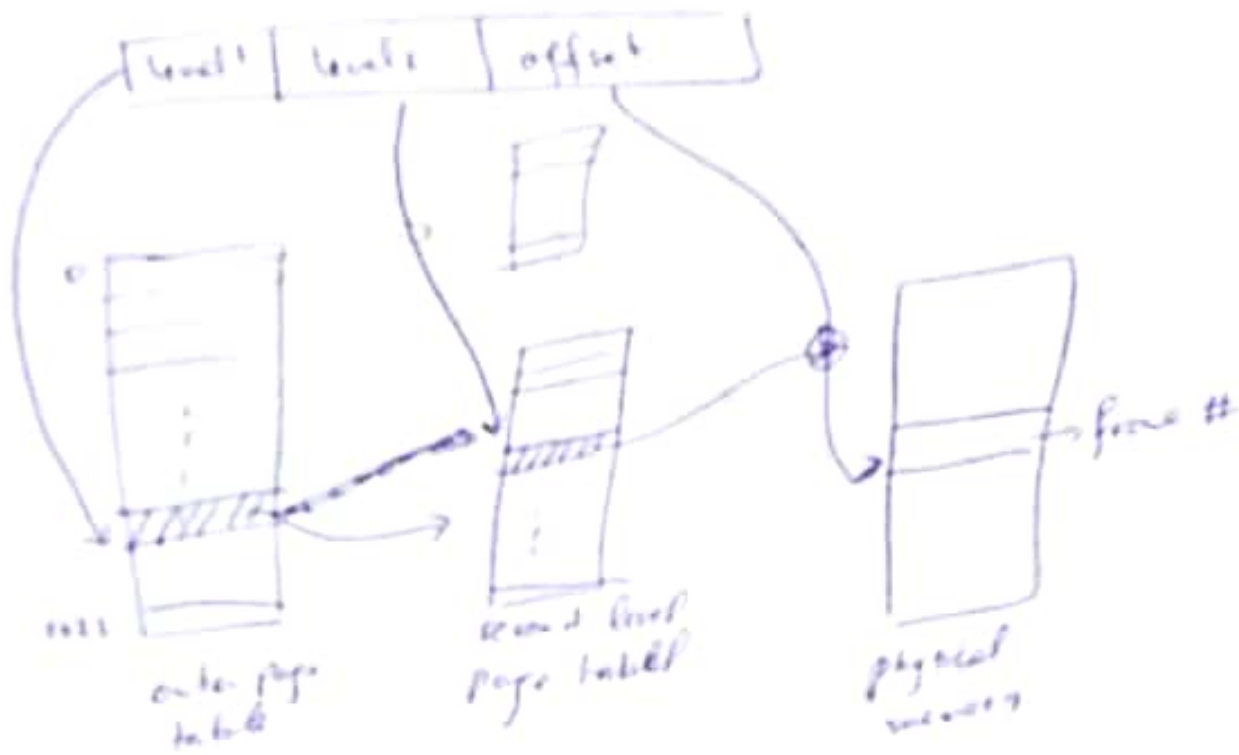
```
    while (i <= 100) sleep(10);
```

```
    else exit(1);
```

```
  } return 0;
```



- the second page index is used to index the second level page table pointed to by the outer page table
- the entry of the second level page table might contain the page frame number of the desired page



For example:  
for a process with 8MB of code  
4MB heap  
2MB stack

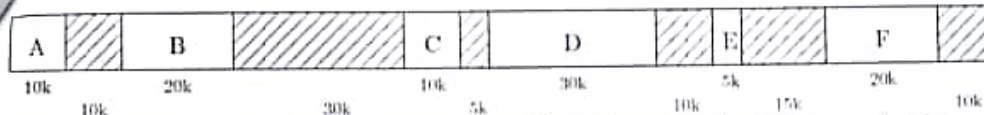
we need ~~3~~ 2 entries for the code  
1 entry for the stack  
1 entry for the heap

So we need just the outer page table  $\rightarrow 4 \times 1024 \text{ bytes} = 4 \text{KB}$   
 - the inner page table for data and heap  
 - the inner page table for stack

$$\Rightarrow 4 \text{KB} + 4 \text{KB} + 4 \text{KB} = 12 \text{KB}$$

while in single-level page table it needs 4MB

to the process. At the instant  $t$ , the state of the main memory is described in the figure below (where A, B, C, D, E, F are processes already loaded in memory, and the other parts are free zones):



Represent the evolution of the main memory according to the arrival of each of the following successive events.

- Event 1. Arrival of program G (20k).
- Event 2. Start of program B.
- Event 3. Arrival of program H (15k).
- Event 4. Departure of the program E.
- Event 5. Arrival of program I (40k).

2. The memory of a computer contains 4-page frames and, at the beginning, all the frames are empty. How many page faults produces the following page references 3, 4, 4, 1, 5, 2, 3, 1, 4 using, respectively, FIFO, OPTIMAL, and LRU replacement algorithms? Justify your answers by showing the contents of the frames after each reference.

#### Problem IV

20 points

1. Describe (in pseudocode) the steps to be followed by the function **void Move (char \* file\_name, int f1, int f2)** which moves a file from a source directory (f1 is the index of its descriptor entry) to a destination directory (f2 is the index of its descriptor entry).
2. Describe (in pseudocode) the steps to be followed by the function **int Similar (int f, char \* file\_name\_1, file\_name\_2)** which compares the contents of two files belonging to the opened directory with entry descriptor f. If both files have the same content, the function returns 1, and 0 if not.

P.S: Feel free to use the structures and functions seen in the course.

Bon travail

## Part 3: VLSI system

### A) indirect list and indirect addressing strategy



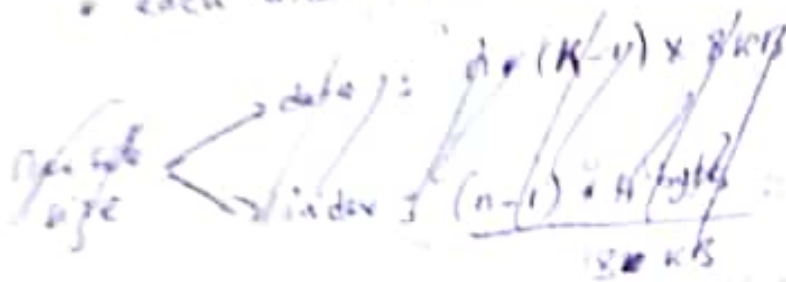
• size block = 8KBS

• pointer occupy 32 bits

Question: Max size (in blocks) of a file?

0 → n-1 blocks ⇒ n blocks

• each block can contains  $8KBS / 4B = 2K \text{ entries}$



• each index block points to  $2K-1$  data blocks = 2047

• pointer occupies 32 bits ⇒  $2^{32}$  blocks on the system

• if  $m$  is the max number of index blocks and all these blocks are used by the file (index + data) then we have?

$$2^{32} = \underbrace{m}_{\text{index}} + \underbrace{(2047 \times m)}_{\text{data}} \Rightarrow m = \frac{2^{32}}{2048} \Rightarrow \boxed{m = 2^{21}}$$

⇒ the number of data block =  $2^{21} \times 2047$   
 ⇒ Max file size is  $2^{32}$  blocks  $\begin{cases} 2^{21} \text{ index block} \\ 2^{21} \times 2047 \text{ data blocks} \end{cases}$

(4)



also // the block to delete is inside a map block but not  
-1...0 shift all entries

B) void delete\_block(int F, int log-b-d)

{ ~~// if the file is short, so~~

// the file is long, so we should know in which map

// block is the block to delete.

// if it is in the last so no problem, we load the

// map block and shift all entries back one

// if not  $\Rightarrow$  we should shift all entries in the map

// block containing the block to delete and all

// other following map blocks

int nb-of-map-blocks=0;  
int map-of-block-to-delete; int remaining-map-blocks=0;

nb-of-map-blocks = fdex[F].lg % 512 == 0 ? fdex[F].lg / 512  
: (fdex[F].lg / 512) + 1;

map-of-block-to-delete = log-b-d / 512;

if (map-of-block-to-delete == (nb-of-map-blocks) - 1)

{ // the block to delete is on the last map block of the file  
// so, just shift the entries of this map block

load(F, map-of-block-to-delete);

x = log-b-d % 512; fdex[F].pbd = fdex[F].map[x];

for (int i = x + 1; i < 512; i++)

map[i-1] = map[i];

block-release(fdex[F].pbd);

}

15 pts

```

else
{
    // the block to delete is inside a map block but not
    // the last one, so we should shift all entries
    // of the following map block one block
    remaining-map-blocks = nb-of-map-blocks - (map-of-
        block-to-delete);
    x = log-2-d % 512;
    for (i = map-of-block-to-delete; i < nb-of-map-block; i++)
    {
        for (j = x+1; j < ((next-map-block * 512) + 512 - x); j++)
        {
            if (j % 512 == 0) load-map(F, j);
            map[j] = map[j-1];
            map[j-1] = map[j % 512];
        } // end for
    } // end for
} // end else
} // end function

```

# 13303 / INFO324 Operating System II

## Problem I

18 points

In each of the following programs, it is assumed that the parent process has  $Pid = 100$ , and that there are no other active processes in the system than those created by the program. Give all possible display results obtained by running each of the following programs.

### Program 1

```
void main() {
    fork();
    if (fork())
        printf("%d\n", fork());
}
```

(5)

### Program 2

```
void main() {
    int x;
    x = (fork() * fork() * fork());
    printf("%d\n", x);
}
```

(2)

### Program 3

```
void main() {
    int p[2], x = 0, y;
    if (fork()) {
        pipe(p);
        x = getpid();
        write(p[1], &x, sizeof(int));
    }
    else {
        read(p[0], &y, sizeof(int));
        printf("%d %d", x, y);
    }
}
```

(5)

## Problem II

16 points

- Using the SIGALRM signal, write a program that draws a number between 1 and 100 after one second and displays it.
- Write a C program under UNIX where the parent process creates 100 child processes as follows: The father creates the first process and waits for the child to display its  $Pid$ , before creating the second child. Then he waits for the second to display his  $Pid$  before creating the third child, and so on.

## Problem III

16 points

- We consider a memory managed by contiguous allocation. The allocation of processes is made according to the first fit algorithm, that is, the first encountered area whose size is greater than or equal to the size of the process to be loaded is the one that is allocated

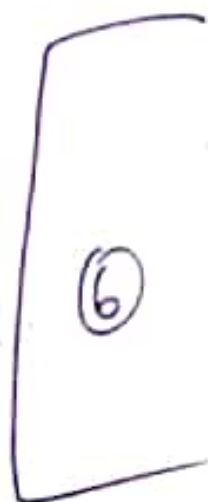


## Problem II :

a)

```

void handler(int sig) { } ②
void main() {
    srand(time(0));
    ① signal(SIGALRM, handler);
    ① alarm(1);
    ① pause();
    ① int num = (rand() % 100) + 1;
    ① printf("%d\n", num);
}
    
```



b)

```

void handler(int sig) {
    printf("I'm the parent and I got a message from my child!\n");
    ③ signal(SIGUSR1, handler);
}
    
```

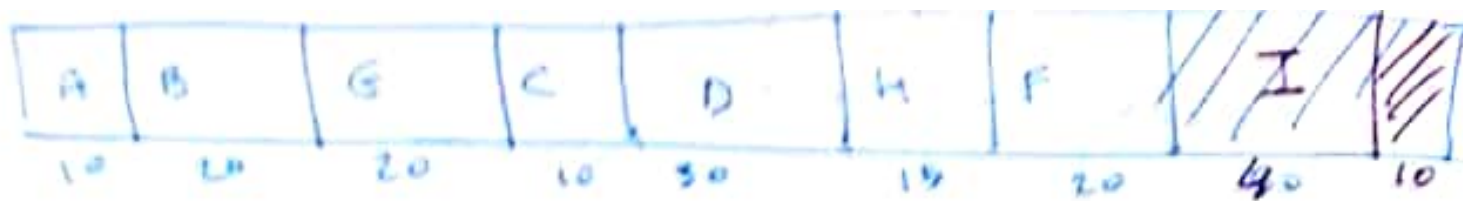
```

void main() {
    ① signal(SIGUSR1, handler);
    ① for(int i=0; i<100; i++) {
        if(fork() pause(); ②
    }
    else {
        ③ printf("I'm the child with rank %d\n", i);
        kill(SIGUSR1, getpid()); ① break;
    } // end if ①
} // End for
} // End main
    
```



and pid = %d\n  
(getpid())

①



2)  $w = \{5, 4, 4, 1, 5, 2, 3, 1, 4\}$

FIFO

(1)

ref	Fault	A	B	C	D
3	Y	3			
4	Y	3	4		
4	N	3	4		
1	Y	3	4	1	
5	Y	3	4	1	5
2	Y	2	4	1	5
3	Y	2	3	1	5
1	N	2	3	1	5
4	Y	2	3	4	5

7 page faults

LRU

(2)

ref	Fault	A	B	C	D
3	Y	3			
4	Y	4	3		
4	N	4	3		
1	Y	1	4	3	
5	Y	5	1	4	3
2	Y	2	5	1	4
3	Y	3	2	5	1
1	N	1	3	2	5
4	Y	4	1	3	2

7 page faults

optimal

optimal

(3)

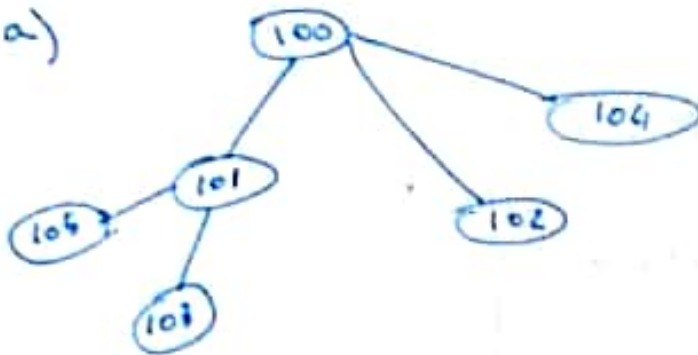
3	4	4	1	5	2	3	1	4
3	3	3	3	3	3	3	3	4
	4	4	4	4	2	2	2	2
			1	1	1	1	1	1
			5	5	5	5	5	5

(2)

6 page faults

problem 1

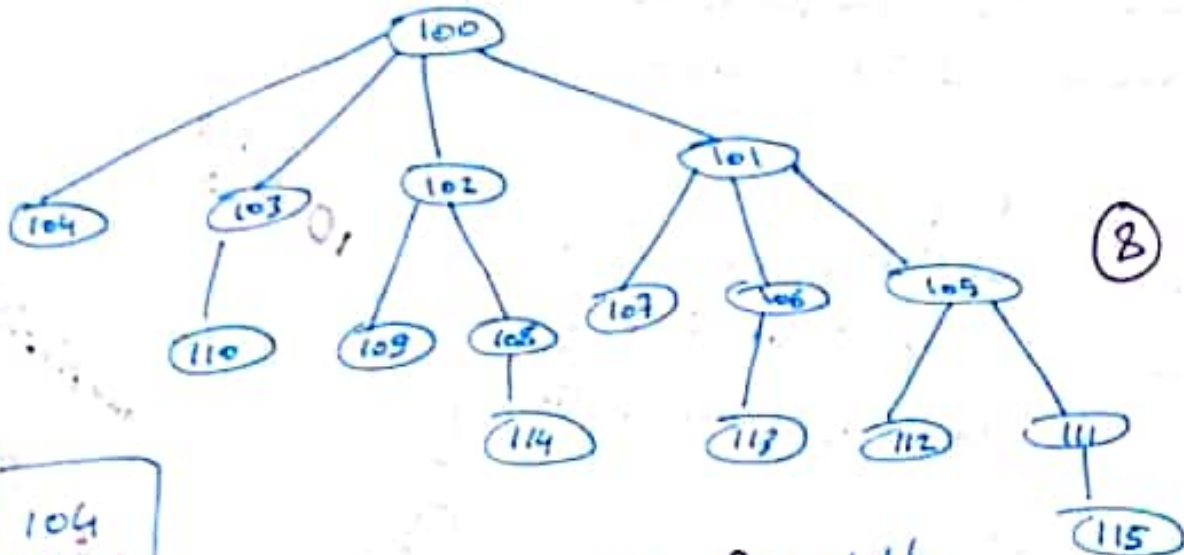
a)



(5)

outputs: 104 } pid of last 2 childs created  
105

b)



(8)

104  
110  
109  
112  
113  
114  
115  
107

8 outputs

other 8 outputs  
prints 0  
(child)

c) the output is 0 0 (5)

(1)



**Part 1 (20 minutes - 15 points)**

A) Answer the following questions with justification :

- What is the advantage of contiguous file allocation policy vs linked allocation policy? (2)
- Describe in detail the problem of fragmentation and how to manage it? (4)

B) Given the following program

```
#include <stdio.h>
#include <unistd.h>
int main() {
    int i;
    for (i=0; i<3; i++)
        if (fork()) i++;
    while(1);
    return i;
}
```

How many processes does this program generate? Draw the generated graph. (5)

C) How many processes are generated by the following code :

```
int main () {
    while (fork())
        execv (path, com);
    return 0;
}
```

Where path is the path to the executable and com is the executable process. (4)

**Part 2: Memory Management (60 minutes - 30 points)**

A) Consider a contiguous memory system with memory allocated as shown below.

100	200	300	400	500	600	700	800	900	1000
0		1		2		3		4	5

Suppose the following actions occur:

- Process E starts and requests 300 memory units. (1)
- Process A requests 400 more memory units. (1) + 1/2
- Process B exits. (1)
- Process F starts and requests 800 memory units. (1)
- Process C exits. (1)
- Process G starts and requests 900 memory units. (1) + 1/2

- Describe the contents of memory after each action using the first-fit algorithm. (6)
- Describe the contents of memory after each action using the best-fit algorithm. (5)
- For this example, which algorithm is best? (1)

**P.S:** you can compact the memory in case of need such that moving used blocks or free blocks

B) Consider a memory paged system with page size 256 bytes. In this system each process is authorized max 4 frames in main memory. The page table of a process P1 is given in the following table.

page	0	1	2	3	4	5	6	7
frame	011	001	000	010	100	111	101	110
presence	1	0	1	0	0	0	1	0

Q. II (20 pts)

void move.(char \*file\_name, int f1, int f2)

folder-entry ent;  
inode binode[63]  
int fd, inode-file;

fd = file-open(file\_name)

~~fd = file-open~~

// f1 and f2 are already linked into  
fd in memory

// Just we should save the link of  
inode of the file for directory f1  
to f2 in memory, then write to disk

// search in data of ~~fd~~<sup>data</sup>[f1] for the  
entry containing file\_name and inode-file  
and delete this entry

// f-inode-file = fd.inode-~~nb~~<sup>nb</sup>

// add new entry in the ~~fd~~<sup>data</sup>[f2]  
write(f2, file\_name);  
write(f2, inode-~~nb~~<sup>nb</sup>);

}

10

1) What is the size of the virtual address space of the process P1 (1)

2) What is the size of the physical memory? (2)

3) Convert the following virtual addresses to physical (signal the error if any): 546, 2072 (0)

4) What will happen if P1 generates the virtual address 770? (1)

C) Consider a system with memory capacity 2GB, page size 4KB and addressing on 32 bits. Given that each table entry contains a reference to a frame + 1 bit presence/absence.

a. What is the size of the page table (justify your answer)

b. How many pages are needed to load the page table in memory?

D) Consider a process with virtual address space of 600 Bytes. the set of virtual addresses referenced is: 34, 123, 145, 510, 456, 345, 412, 10, 14, 12, 234, 336, 412.

a. Give the list of referenced pages given the size of the page is 100 Bytes.

b. Determine the number of page faults for the LRU algorithm. The memory is initially empty and contains 3 frames

#### Part 4: File System (40 minutes – 25 points)

A) Consider a file currently consisting of 100 blocks of data. Assume that the file control block is loaded in memory and there is no cache disk. The size of the block is 4KB. Calculate the number of disk I/O operations required for contiguous and linked allocation strategies to make the following changes to the file. In the contiguous case, you may assume there is no space to grow in the end.

Also assume that the new information to be added to the file is not stored in memory.

a. Add 2 blocks at the beginning (6)

b. Add 2 blocks at the end (6)

c. Remove the middle block (6)

B) Refer to the functions written in class (i.e., `create_inode`, ...):

a. Describe (without writing code) the steps needed to create an inode. (3)

b. How many I/O disk request is required to perform this task. (3)

C) Consider a disk of size 20GB in which the system installed is 16-bit DOS (FAT). The disk is divided into a set of blocks of fixed size (128KB). This disk contains 520 files: 200 files of size 16K, 200 files of size 256KB and 120 files of size 1Mb.

a. Calculate in MB the disk space. (1)

b. Calculate the number of blocks on disk. (1)

c. How many blocks do occupy each of these three categories of files? (2)

d. Calculate in Kb the size of the FAT table. Justify your answer. (3)



```

similar (int f, char *f1, char *f2)
// f is the descriptor of directory loaded into memory
// fdesc[f].topo contains the data block of f
int fd1 = open_file(f1);
int fd2 = open_file(f2);
char *c1, *c2;
if fdesc[f].log <> fd1
if fdesc[fd1].log <> fdesc[fd2].log
    return 0;
else
{
    while (!feof(fd1, c1) != !feof(fd1))
    {
        read(fd1, c1);
        read(fd2, c2);
        if c1 != c2
        if (!strcmp(c1, c2))
            return 0;
    }
    return 1;
}
}

```

(3)

Part I

A) a) advantage of contiguous file allocation vs linked allocation policy

- quick and easy calculation of block holding data & just offset from start of file
- for sequential access, no seeks required
- the read performance is excellent bcz the entire file can be read from the disk in a single operation only one seek is needed (the first block).
- No problem of reliability whereas this is a big problem in linked allocation
- the amount of storage is a power of 2

(2)

b) Problem of fragmentation

During its lifespan, a process can request and free many chunks of memory.

• when a process is started, the free memory areas are long and contiguous.

• Over time and with use, the long contiguous regions become fragmented into smaller and smaller contiguous areas. Eventually, it may become impossible for the program to obtain large contiguous chunks of memory.

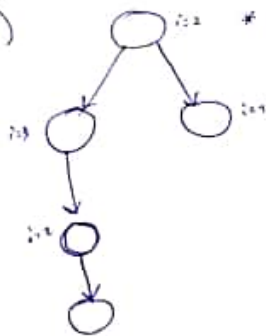
(4)

• there exist two types of fragmentation:

- internal: Due to the rules governing memory allocation (such as paging), more computer memory is sometimes allocated ~~external~~ than is needed. For example, in a system with page size 512 bytes, a file with size 400 KB is allocated one page of 512 KB  $\Rightarrow$  this waste is called internal frag-

## Part I (continued)

B)



\* There is 5 processes including the main process

(5)

C) just one process is generated in this code by the parent (main).



The main process change the address space and executes another process and the child exits.

(4)

## Part II (Memory management)

A)

0	A
400	E
700	
800	
	B
1600	
1900	
2400	C
3100	
3400	D

a) First Fit algorithm:

① \* process E starts and requests 300 memory units  
⇒ allocation in the first free zone 400-700

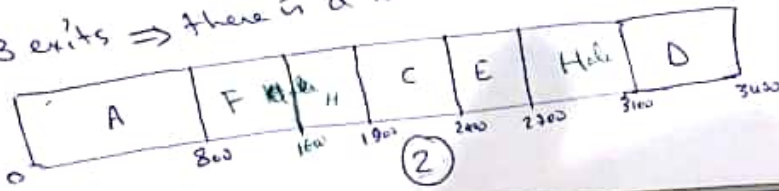
\* process A requests 400 more memory units  
⇒ no free spaces to give to process A

② ⇒ compact the memory  
⇒ cannot fit because the entire process is allocated in a single continuous chunk of memory in a contiguous memory system

⇒ Move B to 1100-1900  
Move E to 2400-2700

\* Give A additional addresses 400-800

\* B exits ⇒ there is a hole between 800-1900



(2)



- External : arises when free memory is separated into small blocks and is interspersed by allocated memory.
- The term external refers to the fact that the unusable storage is outside the allocated regions.
- For example in a contiguous allocation strategy with ~~dynamic~~ partitions, consider a situation where in a ~~page~~ <sup>the memory allocator</sup> allocates 3 contiguous ~~blocks~~ regions of memory for 3 different processes, and then free the middle region due to a swap or exit. if a new demanded region of memory is larger than the free space  $\Rightarrow$  this free space is called external fragmentation.

$\rightarrow$  Internal fragmentation : view by the process  
 $\rightarrow$  external frag-ent : view by the system

### Remediation

- there is no complete solution for internal fragmentation.
- the internal fragmentation is always, at max with the size of a memory page.
- \* the external fragmentation can be avoided by ~~using~~ decomposing the memory into fixed size blocks such as paging system.

- \* process F starts and requests 800 memory units  
 $\Rightarrow$  F is allocated in 800-1600 (1)
- \* C exists  $\Rightarrow$  there is a hole between 1600-2400 (1/2)
- \* G starts and requests 900 memory units  
 $\Rightarrow$  no hole that is big enough  
 $\Rightarrow$  compact memory: move E to 2800-3100 (1) + (1/2)  
 $\Rightarrow$  G is allocated in 1600-2500

### b) Best-Fit algorithm

- E requests 300  $\Rightarrow$  E is allocated in 1600-1900 (1)
- A requests 400 more  $\Rightarrow$  this 400 is allocated in 400-800 (1)
- B exits  $\Rightarrow$  there is a hole between 800-1600 (1/2)
- F requests 800  $\Rightarrow$  F is allocated in 800-1600 (1)
- C exits  $\Rightarrow$  there is a hole between 1900-3100 (1/2)
- G requests 900  $\Rightarrow$  G is allocated in 1900-2800 (1)

### c) Worst-Fit Algorithm (not required)

- E requests 300  $\Rightarrow$  E is allocated in 2400-2700
- A requests 400 more  $\Rightarrow$  this additional 400 is allocated in 400-800
- B exits  $\Rightarrow$  there is a hole between 800-1900
- F requests 800  $\Rightarrow$  F is allocated in 800-1600
- C exits  $\Rightarrow$  there is a hole between 1600-2400
- G requests 900: no hole big enough  $\Rightarrow$  need to compact  
 $\Rightarrow$  move E to 2800-3100, give 1600-2500 to G.

For this example, Best-Fit is the best (1)

(3)

(2)

C) Memory size = 2 GB  
 Page size = 4 KB  
 addressing 32 bits

each PTE: reference to frame + 1 bit P/A

(a) Size of page table

- addressing on 32 bits  $\Rightarrow 2^{32}$  virtual address space
- page size = 4 KB =  $2^{12}$

$$\Rightarrow \begin{array}{c|c} 20 & 12 \\ \hline \text{page \#} & \text{offset} \end{array}$$

$\Rightarrow$  page table contains  $2^{20}$  entries

Or the number of frames is:  $\frac{\text{Memory size}}{\text{Page size}} = \frac{2 \text{ GB}}{4 \text{ KB}} = \frac{2^{31}}{2^{12}}$

(b)  $= 2^{19}$

$\Rightarrow$  we need 19 bits for the frame #

$\Rightarrow$  each PTE size is  $19 + 1 = 20$  bits

$\Rightarrow$  the size of page table is  $2^{20} \times 20$

$$= 2^{22} \times 5 \text{ Bytes} = 20 \text{ MB} = 20480 \text{ KB}$$

(b) How many pages?

#pages =  $\frac{20480 \text{ KB}}{4 \text{ KB}} = \frac{20480}{4} = 5120$  pages

D) Process: virtual address space 600 Bytes page size = 100 Bytes

a) 0, 1, 1, 5, 4, 3, 4, 0, 0, 0, 2, 3, 4

9 page faults

Page #	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0	0	0	0	4	4	4	4	4	4	4	3	3
1		1	1	1	1	3	3	3	3	3	2	2	2
2				5	5	5	5	0	0	0	0	0	4

4



B) - page size = 256 Bytes

- each process is authorized max 4 frames

Page	0	1	2	3	4	5	6	7
Frame	011	001	000	010	100	111	101	110
Process	1	0	1	0	0	0	1	0

1) Size of virtual address space of P1

$$8 \times 256 = 2048 \text{ bytes } \textcircled{2}$$

2) Size of physical memory

The frames are coded in 3 bits  $\Rightarrow 2^3 = 8$  frames  
 $\Rightarrow$  physical memory size = nb(frames)  $\times$  size(frame)  
 $\textcircled{2} = 8 \times 256 = 2048 \text{ bytes}$   
 $= 2 \text{ KB}$

3) convert to physical address

The ~~conversion~~ conversion from virtual address to physical is realized as the following:

- calculation of page # and offset
- search in page table the frame #
- physical address ~~frame~~ frame address + offset

then:

$$546 = 2 \times 256 + 34$$

$\Rightarrow$  page 2 & offset 34

$\Rightarrow$  frame 0 (page table)

$\Rightarrow$  physical address is 34  $\textcircled{1}$

2072 is out of the virtual address space of P1  
 $\Rightarrow$  error  $\textcircled{1}$

4)  $770 = 3 \times 256 + 2 \Rightarrow$  page 3. But this page is not loaded in memory  $\Rightarrow$  page fault  $\textcircled{1}$

$\textcircled{3}$

## Part 4: File System

- file: 100 blocks of data
  - File control block is loaded in memory
  - size of block = 4KB
- Number of disk I/O?

	contiguous	linked
Add 2 blocks at begin	$2R + 2W$ (1)	$2R + 2W$ (1)
Add 2 blocks at end	$102R + 102W$ (2)	$3R + 3W$ (2)
Remove the middle block	$50R + 50W$ (1)	$\cancel{2R} + 1W$ (2) <del>50</del>

i) add 2 blocks at the beginning

- Contiguous: 1 reads block for new information + 2 write blocks "
- linked: 2 reads for new inf
  - update pointers in memory
  - 2 writes for new inf

ii) add 2 blocks at the end

- contiguous: there is no room to add at the end  
 $\Rightarrow$  shifting all the blocks two places at the end

$$\Rightarrow 100R + 100W$$

- read 2 new inf blocks
- write 2 new inf blocks }  $2R + 2W$

$$\Rightarrow \boxed{102R + 102W}$$

- linked:
    - Read in the last block  $\rightarrow 1R$
    - Read in the 2 new blocks  $\rightarrow 2R$
    - update pointers in memory  $\rightarrow 0R, 0R$
    - write the (4) three blocks  $\rightarrow 3W$
- $\left. \begin{array}{l} 3R \\ + \\ 3W \end{array} \right\}$

ii) Remove the middle block

- contiguous:

to delete position 50, read all blocks <sup>after</sup> ~~before~~ 50  
and write them back one place ~~closer to the front~~  
 $\Rightarrow 50R + 50W$

- ~~read 2 new inf blocks and write it back~~  
 $\Rightarrow \text{2R + 2W}$

$$\Rightarrow \boxed{50R + 50W}$$

- linked:

read and follow links to position 50  
 $\Rightarrow 50R \rightarrow$  update in memory block 49 (Next)  $\rightarrow 51$   
then write block 49 to link to former  
block 51  $\Rightarrow 1W$

$$\Rightarrow \boxed{50R + 1W}$$

B) create - inode

a) the steps to create inode are:

- search for free inode:

1- in cache

2- not found, in disk

(2) (3)

- bring the block that contains the inode to ~~any~~

- initialize the inode

- write back the block containing the inode to disk.

b) # of disk I/O?

\* if there is free inode in cache  $\Rightarrow$  ~~cache~~  
 $1R + 1W$

\* if no free inode in cache  $\Rightarrow$  search on disk

$\Rightarrow$  worst case read all blocks of inodes

$$\Rightarrow \boxed{size R + 1W}$$

(2) (3)