

Qb1 (Part A)

void main ()

{ fork(); /* A */

if (fork())

printf ("%d\n", fork());

}


1004

1005

0

0

Part B: (25 pts)



1) (8 pts)

#define N 10

int main ()

{ int found = 0;

int pid[N];

int current, int parentpid;

if (fork == 0)

break;

// process Q

else // Parent context

{ for (i=1; i<N; i++)

if (pid[i] == fork())

break;

}

sleep(2);

current = getpid();

for (i=1; i<N; i++)

{ if (current == pid[i])

found = 1;

```

for (j=1; j ≤ n; j++)
    if (fork() == 0)
        break;
}

```

return 0;

} 2) (7 pts)


2) parentpid = getpid()

① /* we put this statement before creating all other processes */

② if (parentpid == getpid())
fork(); /* creation of 0 process */

③ while (wait(NULL));
/* put in the loop */

(6)



Page tables:

P_1

Page #	Frame #	Pres	R
0	1	1	1
1	9	1	1
2	-	0	0
3	6	1	1

P_2

Page #	Frame #	Pres	Ref
0	3	1	1
1	7	1	1
2	8	1	1
3	-	0	0
4	-	0	0

Page #	Frame #	Pres	Ref
0	0	1	1
1	-	0	0
2	4	1	1
3	5	1	1

< Page 12, offset 4 >


: access out of memory =>
 ① segmentation fault

- 2) $P_3 \equiv \langle \text{page 2, offset 14} \rangle \equiv (3, 2) \Rightarrow \text{hit}$ (2)
- 3) $(P_3, \text{page 2}) \equiv (3, 2) \Rightarrow \text{hit}$
- $(P_1, \text{page 2}) \equiv (1, 2) \Rightarrow \text{fault page} \Rightarrow \text{load into frame 2}$
- $(P_3, \text{page 2}) \equiv (3, 2) \Rightarrow \text{hit}$
- $(P_2, \text{page 2}) \equiv (2, 2) \Rightarrow \text{hit}$
- $(P_2, \text{page 4}) \equiv (2, 4) \Rightarrow \text{fault page} \Rightarrow \text{load into frame 9}$
- $(P_2, \text{page 3}) \equiv (2, 3) \Rightarrow \text{fault page} \Rightarrow \text{load into frame 6}$
- $(P_1, \text{page 0}) \equiv (1, 0) \Rightarrow \text{fault page} \Rightarrow \text{load into frame 3}$
- $(P_3, \text{page 3}) \equiv (3, 3) \Rightarrow \text{hit}$
- $(P_1, \text{page 1}) \equiv (1, 1) \Rightarrow \text{fault page} \Rightarrow \text{load into frame 7}$
- $(P_1, \text{page 2}) \equiv (1, 2) \Rightarrow \text{hit}$
- \Rightarrow The no. of page faults is $\frac{9}{2} + 5 = 14$

12 entries

block size = 16 KB

block # = 2 bytes



1) what is the max size of a file?

16 KB = $16 \times 1024 = 2^4 \times 2^{10} = 2^{14}$ bytes (3)

\Rightarrow the block occupy $2^{14} \div 2 = 2^{13} = 8192$ entries

\Rightarrow the max size is : $10 \times 16 + 8191 \times 16 + 8192 \times 8192 \times 16 =$
 $10 \times 2^4 + 2^{13} \times 2^4 + 2^{13} \times 2^{13} \times 2^4 =$
 $10 \times 2^4 + 2^{17} + 2^{30}$ Kbytes

a) $x = 12$

$y = 8192$

$z = 8192$

$t = 2^4 \times 2^{10} = 2^{14}$ (4)

b) data-free (int f, in ldb)

(release the logical data block of from the file f)

(2)

```

void free-data (int F, int logbd)
{
    fdesc[F].pbd = fdesc[F].topo[logbd];
    block-release (fdesc[F].pbd);
}
if (! fdesc[F].pbd)
    fdesc[F].topo[logbd] = 1;
}

```

c) u void map1-free (int F, int logbm)

```

{
    load-map1 (F, logbm);
    // the buffer map1 contains the map block
    for (i=0; i < 8192; i++)
        block-release (fdesc[F].map1[i]);
    block-release (fdesc[F].topo[logbm]);
}

```

d) u void map2-free (int F, int logbm)

```

{
    load-map2 (F, logbm);
    // load into buffer map2
    for (i=0; i < 8192; i++)
    {
        load-map2 (F, fdesc[F].map1[i]);
        for (j=0; j < 8192; j++)
        block-release (fdesc[F].map2[j]);
        map1-free (fdesc[F].map2[i]);
    }
}

```

void file_delete (int F)
{

~~free_data(F);~~

for (i=0; i<10; i++)

free_data (~~free_data~~ F, i);

map1_free (F, 11);

map2_free (F, 12);

inode_release (fdesc[F].inode_nb);

}

6

Process management


a) using 2 or more level page table

since each page in the virtual address space needs an entry in the page table, we may have a huge number of entries in the page table when actually a few of them are needed by the process. for example in a virtual address space of 32 bits with page of size 4KB = 2^{12} we need 2^{20} entries. Each entry occupy 4 bytes

\Rightarrow 4KB the page table are reserved for a process which it actually needs a few entries. For a 64 bits virtual address space we need $2^{52} \times 4$ bytes for the page table for each process

\rightarrow using multi-level page table can reduce the overhead as follows:

for example in two-level page table the virtual address space is divided into 3 parts: (32 bits)




(2)

there are two page table indexes:

- the first page index is used to index the outer page table
- each entry in the outer page table can possibly point to a second level page table
- the present bit in the outer page table is used to indicate if there is a second level page table for this entry

(1)

- * the second page index is used to index the second level page table pointed to by the outer page table
- * the entry of the second level page table might contain the page number of the desired page



For example:

for a process with 8MB of code
4MB heap
2MB stack

we need ~~2~~ 2 entries for the code
1 entry for the stack
1 entry for the heap

So we need just the outer page table $\rightarrow 4 \times 1024 \text{ bytes}$
 $= 4 \text{ KB}$


- the inner page table for data and heap
- the inner page table for stack

$$\Rightarrow 4 \text{ KB} + 4 \text{ KB} + 4 \text{ KB} = 12 \text{ KB}$$

while in single-level page table it needs 4 MB

or a physical address

Logical address is used by each process in ~~each~~ its virtual address space starting from 0 to max-1 bytes depending on the architecture and it is private for each process. whereas the physical address space is the real address on physical memory. It has two parts



B) a) zombie state

a zombie process is a process that has completed execution (via the exit system call) but still has an entry in the process table waiting his parent to get its status by the (wait call) (2)

b) orphan state

~~process~~ An orphan process is a process that is still executing, but whose parent has died (2)

→ #include <unistd.h>
#include <stdio.h>

int main()

{ int pid; int i=0; (5)

pid = fork();

if (!pid)

while (i <= 100) sleep(10);

else exit(1);

} return 0; (2)


Memory management (15 pts)

4 pages

→ 2 pages

size of entry:

$$64 \times 4 = 256 \text{ Bytes} = 2^8 \text{ Bytes} \quad (3)$$



2) size of the virtual address space

~~the~~ Page size = 64 By = 2^6 Bytes

⇒ 6 bits $\hat{=}$ needed for the offset

- the logical address is in hexadecimal (2 digits)

- each is encoded on 4 bits ⇒ the logical address is encoded on 8 bits ⇒ $\frac{4}{2} \times \frac{4}{6}$

$$\Rightarrow 2^8 \text{ pages} \Rightarrow 4 \times 64 = 256 \text{ Bytes} = 2^8 \quad (3)$$

3F	⇒	0011 1111	63	0	P ₁
4A	⇒	0100 1010	74	1	P ₁
1D	⇒	0001 1101	29	0	P ₁
00	⇒	0000 0000	0	0	P ₂
CA	⇒	1100 1010	202	3	P ₁
87	⇒	1000 0111	135	2	P ₁
39	⇒	0011 1001	57	0	P ₂
2B	⇒	0010 1011	43	0	P ₁
00	⇒	0000 0000	00	0	P ₁
11	⇒	0001 0001	17	0	P ₂

⇒ the $W = \{0, 1, 0, 0, 3, 2, 0, 0, 0, 0\}$

$$W = \{(0; P_1), (1; P_1), (0; P_1), (0; P_2), (3; P_1), (2; P_1), (0; P_2), (0; P_1), (0; P_1), (0; P_2)\}$$

1) FIFO

	0	0	0	0	0	0	0	0	0
1		1	1		1	1	1	1	1
2					3	3	3	3	3
3					2	2	2	2	2

4 page faults

a) FLFC

Page refs	4 page frames				
	Fault?	Page contents			
(0, P ₁)	✓	(0, P ₁)			
(1, P ₁)	✓	(1, P ₁)	(0, P ₁)		
(0, P ₁)	X	(1, P ₁)	(0, P ₁)		
(0, P ₂)	✓	(0, P ₂)	(1, P ₁)	(0, P ₁)	
(3, P ₁)	✓	(3, P ₁)	(0, P ₂)	(1, P ₁)	(0, P ₁)
(2, P ₁)	✓	(2, P ₁)	(3, P ₁)	(0, P ₂)	(1, P ₁)
(0, P ₂)	X	(2, P ₁)	(3, P ₁)	(0, P ₂)	(1, P ₁)
(0, P ₁)	✓	(0, P ₁)	(2, P ₁)	(3, P ₁)	(0, P ₂)
(0, P ₁)	X	-	-	-	-
(0, P ₂)	X	-	-	-	-

6 page faults (3)

Page refs	Fault?	4 Page Frames			
		Page contents			
(0, P ₁)	✓	(0, P ₁)			
(1, P ₁)	✓	(1, P ₁)	(0, P ₁)		
(0, P ₁)	X	(0, P ₁)	(1, P ₁)		
(0, P ₂)	✓	(0, P ₂)	(0, P ₁)	(1, P ₁)	
(3, P ₁)	✓	(3, P ₁)	(0, P ₂)	(0, P ₁)	(1, P ₁)
(2, P ₁)	✓	(2, P ₁)	(3, P ₁)	(0, P ₂)	(0, P ₁)
(0, P ₂)	X	(0, P ₂)	(2, P ₁)	(3, P ₁)	(0, P ₁)
(0, P ₁)	X	(0, P ₁)	(0, P ₂)	(2, P ₁)	(3, P ₁)
(0, P ₁)	X	-	-	-	-
(0, P ₂)	X	(0, P ₂)	(0, P ₁)	(2, P ₁)	(3, P ₁)

5 page faults (3)

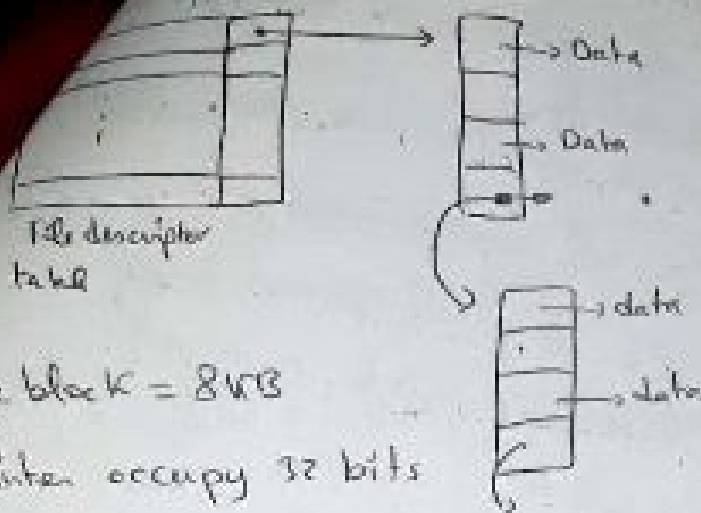
c) Second chance

Page ref	Fault?	4 Page Frames					out
		Page contents					
(0, P ₁)	✓	(0, P ₁)					
(1, P ₁)	✓	(1, P ₁)	(0, P ₁)				
(0, P ₁)	X	(1, P ₁)	(0, P ₁) [*]				
(0, P ₂)	✓	(0, P ₂)	(1, P ₁)	(0, P ₁) [*]			
(3, P ₁)	✓	(3, P ₁)	(0, P ₂)	(1, P ₁)	(0, P ₁) [*]		
(2, P ₁)	✓	(2, P ₁)	(3, P ₁)	(0, P ₂)	(0, P ₁)	(0, P ₁)	
(0, P ₂)	X	(2, P ₁)	(3, P ₁)	(0, P ₂) [*]	(0, P ₁)		
(0, P ₁)	X	(2, P ₁)	(3, P ₁)	(0, P ₁) [*]	(0, P ₁) [*]		
(0, P ₁)	X	—	—	—	—		
(0, P ₁)	X	—	—	—	—		

5 page faults

(3)

indirect list and indexed allocation strategy



size block = 8KB

pointer occupy 32 bits

(10)

Question : Max size (# blocks) of a file ?

0 → n-1 blocks → n blocks

* each block can contains $8KB / 4B = 2K$ entries.

$$\text{File size} = \text{data} + \text{index} = n \times (K-1) \times 8KB + (n-1) \times 4 \text{ bytes}$$

$32KB$

* each index block points to $2K-1$ data blocks = 2047

* pointer occupies 32 bits $\Rightarrow 2^{32}$ blocks on the system

* if m is the max number of index blocks and all these blocks are used by the file (index + data) then we have:

$$2^{32} = \underbrace{m}_{\text{index}} + \underbrace{(2047 \times m)}_{\text{data}} \Rightarrow m = \frac{2^{32}}{2048} \Rightarrow \boxed{m = 2^{21}}$$

\Rightarrow the number of data block = $2^{21} \times 2047$

\Rightarrow Max file size is 2^{32} blocks

\swarrow 2^{21} index block
 \searrow $2^{21} \times 2047$ data blocks

(4)

B) void delete_block (int r, int log-u-u,

{ ~~// if the file is short, so~~

// the file is long, so we should know in which

// block is the block to delete.

// if it is in the last so no problem, we load the

// map block and shift all entries back one

// if not \Rightarrow we should shift all entries in the map

// block containing the block to delete and all

// other following map blocks

int nb_of_map_blocks = 0;

int map_of_block_to_delete; int remaining_map_blocks = 0;

nb_of_map_blocks = fdesc[F].lg / 512 == 0? fdesc[F].lg / 512
: (fdesc[F].lg / 512) + 1;

map_of_block_to_delete = log-b-d / 512;

if (map_of_block_to_delete == (nb_of_map_blocks) - 1)

{ // the block to delete is on the last map block of the file

// so, just shift the entries of this map block

load(F, map_of_block_to_delete);

x = log-b-d % 512; fdesc[F].pbd = fdesc[F].map[x];

for (int i = x + 1; i < 512; i++)

map[i-1] = map[i];

block_release(fdesc[F].pbd);

15 pts

// the last one, so we should shift all entries
 // of the following map block one back

remaining-map-blocks = nb-of-map-blocks - (map-of-

$x = \log_2 d \% 512;$ block-to-delete);

for (i = map-of-block-to-delete; i < nb-of-map-block; i++)

{
 for (j = x + 1; j < ((new-map-block * 512) + 512 - x); j++)

{
 if (j % 512 == 0) load-map(F, j);

~~map[j] = map[j]~~

map[(j-1) % 512] = map[j % 512];


} // end for

} // end for

} end else

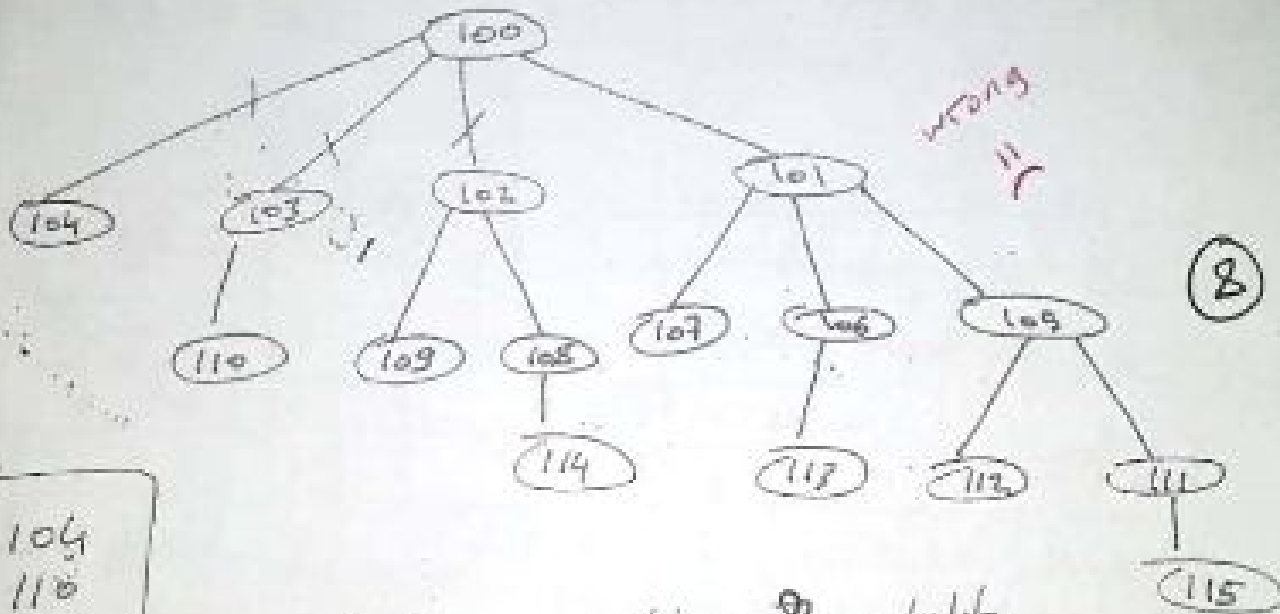
} // end function

problem 1



(5)

output: 104 } pid of last 2 childs created
105



(2)

2 outputs

other 2 outputs
prints 0
(child)

104
110
109
112
113
114
115
107

the output is 0 0 (5)

(1)

a) void handle(int sig) { } ②

void main() {

svrand (time(0));

① signal (SIGALRM, handle);


① alarm(1);

pause();

① int num = (rand() % 100) + 1;

① printf ("%d\n", num);

}



b) void handler (int sig) {

printf ("I'm the parent and I got a message from my child\n");

③ signal (SIGUSR1, handler);

}

void main() {

① signal (SIGUSR1, handler);

① for (int i=0; i<10; i++) {

if (fork()) pause(); ②

else {

③ } printf ("I'm the child with rank %d\n", i);

kill (SIGUSR1, getpid()); break;

} // end if ①


} // End for

} // End main




and pid % 10

getpid()



no place for pg 3 (work) \Rightarrow fragmentation (5)

\Rightarrow



2) $w = \{3, 4, 4, 1, 5, 2, 3, 1, 4\}$

FIFO

LRU

optimal

(1)

ref	Fault	A	B	C	D
3	Y	3			
4	Y	3	4		
4	N	3	4		
1	Y	3	4	1	
5	Y	3	4	1	5
2	Y	2	4	1	5
3	Y	2	3	1	5
1	N	2	3	1	5
4	Y	2	3	4	5

7 page faults

(2)

ref	Fault				
3	Y	3			
4	Y	4	3		
4	N	4	3		
1	Y	1	4	3	
5	Y	5	1	4	3
2	Y	2	5	1	4
3	Y	3	2	5	1
1	N	1	3	2	5
4	Y	4	1	3	2

7 page faults

optimal

(3)

6 page faults

3	4	4	1	5	2	3	1	4
3	3	3	3	3	3	3	3	4
	4	4	4	4	2	2	2	2
			1	1	1	1	1	1
			5	5	5	5	5	5

(2)


```
void move.(char *file_name, int f1, int f2)
{
    folder_entry ent;
    inode binode[16]
    int fd, inode-file;
```

(10)

```
fd = file-open(file_name)
```

```
fd = file-open
```

// f1 and f2 are already loaded into
fdesc in memory

// Just we should move the link of
inode^{f1} of the file for directory f1
to f2 in memory, then write to disk

// search in data of ~~fdesc~~^{desc}[f1] for the
entry containing file name and inode-file
and delete this entry

// ~~inode-file~~ = fd. inode-ub.

// add new entry in the ~~fdesc~~^{desc}[f2]
write (f2, file-name);
write (f2, inode-ub);

}

similar (int f, char *f1, char *f2)

// f is the descriptor of directory loaded into memory

// fdes[f].topo contains the data block of f

int fd1 = open_file(f1);

int fd2 = open_file(f2);

char *c1, *c2;

~~if fd1 <> fd2~~

if fdes[fd1].log <> fdes[fd2].log

return 0;

else

{ while (~~and(f1, c1) || !Eof(fd1)~~)

{

read(fd1, c1);

read(fd2, c2);

~~if c1 != c2~~

if (!strcmp(c1, c2))

return 0;

}

return 1;

}

}

(3)

INFO 324
Operating System II

Dr. Ahmad Fakhour

40 points

Problem I

The parts A et B are independent

Part A)


Consider the following C program under Unix:

```
void main()
{
    fork();
    if (fork())
        printf("%d\n", fork());
}
```

Assume that the parent has an identity $Pid = 1000$, and there is no active processes other those created by the program. Display all the results obtained by executing the above program

Part B)

1. Write a C program under Unix that creates the following tree of processes.



The parent process creates the process P_1 , and each process P_i creates a process P_{i+1} ($i=1 \dots N-1$). In addition the parent creates another process Q ; and each process P_i ($i=1 \dots N$) creates a sub-tree of M processes.

Now, we add to the program in part (1) the following two criteria:

- The process Q must be the last process created in the tree (global) of processes.
- After the creation of all processes in the tree, all these process can be in running mode, i.e. the creation of a process should have no relation with the termination of another process

Write the program (the extension of Part 1) that creates the tree of processes in accordance with the two criteria described above.

Problem II

Memory

20 points

Consider a machine where the memory is managed by the mechanism of paging. The physical memory is divided into a set of 10 frames (frames) numbered from 0 to 9. In this context, three processes P_1 , P_2 and P_3 which address spaces are respectively composed of four pages, five pages and four pages (the pages are also numbered from 0). At time t , the configuration of main memory is as follows:

- For P_1 : pages 0, 1 and 3 are loaded in the frames 1, 9 and 6 respectively.
- For P_2 only pages 0, 1 and 2 are loaded into the frames 3, 7 and 8 respectively.
- For P_3 pages 0, 2 and 3 are loaded in the frames 0, 4 and 5 respectively. The frames not mentioned above are free.

- Represent the page tables of processes according to the previous configuration. In this context, the process P_2 requests access to the virtual address <page 12, offset 4>, what's going on?
- The process P_3 then asks access to the virtual address, <page 2, offset 14>, what's going on?
- We assume that the pages are loaded in that order i.e., page 0 is loaded from P_1 , page 1 and page 3; after page 0 is loaded for P_2 , page 1 and page 2 and likewise for P_3 . What will be the behavior of the

memory for the following requests as follows: (P3, page2), (P1, page 2), (P3, page2), (P2, page2), (P2, page4), (P2, page3) (P1, page0), (P3, page3), (P1, page 1), (P1, page 2) if we adopt the FIFO strategy? What is the number of page faults?

25

25

points

We consider a FS where the topo table contains 12 entries, each of the first 10 entries point directly to a data block, the 11th entry has only one level of indexing, and the 12th entry has two levels of indexing. The size of a block is 16 KB (kilobytes), and the number of a block occupies two bytes.

1. What is the maximum size of a file in this system?
2. Descriptors files in main memory are assumed open in the following table fdesc:

```
#define maxf
struct {
    int lg;
    int bld, bpd, bic1, bpc1, bic2, bpc2;
    int topo[x];
    int map1[y];
    int map2[z];
    char buffer[t];
    ...
} fdesc(maxf);
```

- Descriptor may be complemented if necessary by other fields covered in course.

- a. Give the values of x, y, z and t, 4
- b. Write the function data_free (int f, int blocd) that permits the release of the logical data bloc (blocd) of the file loaded in the entry f of table of descriptors.
- c. write the function map1_free (int f, int blogm) that permits the release of the logical map bloc of first level of indexing (blogm)
- d. write the function map2_free (int f, int blogm) that permits the release of the logical map bloc of second level of indexing (blogm)
- e. Write the function file_delete (int f) that deletes the file loaded into the entry f, deleting a file is to release all the blocks reserved by the file and its inode.

The use of other functions seen in the course (supposedly adapted to this FS) is freely permitted.

Part 1: Process Management (30 points):

A) Answer the following questions with justification :

- Explain in detail how the use of two-level or more page tables will reduce the space required for the one-level page table. (3)
- Explain the difference between logical and physical addresses (3)

B) After explaining what it means:

- A process is in the zombie state. (2)
- A process is orphaned. (2)

Write a program that creates a process that fall within the zombie state before becoming orphan. (5)

C) Write a C program where a parent creates N child processes. Then, the processes behave in the following manner:

Every T second, the father write in a pipe a character 'G' and sends to all "alive" child the "SIGUSR1" signal for unlock them. The child awakened by the signal, start the race to read the character already written in the tube by the father. The winning child process must terminate after displaying a number indicating the termination order. (5)

P.S: Remember that the father sent the signal "SIGUSR1" only for living child (who are not already terminated.)

Part 2: Memory Management (15 points)

Consider a system with paginated main memory. The memory is composed of 4 frames (frames) each has a size of 64 bytes. At a given moment, the memory is empty, then two processes P1 (4 pages) and P2 (2 pages) are launched in the system. The processor sends requests submissions in the following order in the format [hexadecimal logical address, Process]:

[3F, P1]
[4A, P1]
[1D, P1]
[00, P2]
[CA, P1]
[87, P1]
[39, P2]
[2B, P1]
[00, P1]
[11, P2]

- What is in bytes the size of the main memory? (3)
- What is the size of the virtual address space? (3)
- Indicate by figures the evolution of the memory and the number of page faults using the following page replacement algorithms :
 - FIFO (3)
 - LRU (3)
 - Second chance. (3)

Part 3: File System (25 points)

A) Consider a file system with linked list and indexed allocation strategy as illustrated in the figure. The size of block is 8 KB. $= 8 \times 2^{10} \text{ bytes}$




Figure: File system illustration

The disk blocks are numerated from 0 to $n-1$. Each opened file has an entry in the file descriptor table which is loaded in memory. This entry contains all attributes of the file and a pointer to an index block. The index block contains k pointers to other blocks. These pointers except the last points to data block. The last pointer points to another index block and so on as shown in figure. The pointer to block occupies 32 bits. $\Rightarrow 4 \text{ bytes}$

Question: What is the maximum size (in number of blocks) of a file in this system? Indicate the number of data blocks and index blocks

- Given a FS where the topo table contains 10 entries, each corresponding to a single level indexing (each block points to a map). Given that each block occupies 2 kilobytes, the number of a block occupies 4 bytes and each block contains 16 inode:
Write the function `delete_block(int lbd, int f)` that delete the physical data block corresponding to the logical data block lbd of the file with descriptor f.
Note that if this block is not at the end of file, you must shift all the other data blocks that follow.
P.S: the use of the functions seen in the course is permitted if needed.

18 points

Problem I

In each of the following programs, it is assumed that the parent process has PId = 100, and that there are no other active processes in the system than those created by the program. Give all possible display results obtained by running each of the following programs.

Program 1

```
void main()
{
    fork();
    if(fork())
        printf("%d\n", fork());
}
```


(5)

Program 2

```
void main()
{
    int x;
    x = (fork() + fork()) * fork();
    printf("%d\n", x);
}
```

(2)

fork
fork
fork



Program 3

```
void main()
{
    int p[2], x = 0, y;
    if(fork())
    {
        pipe(p);
        x = getpid();
        write(p[1], &x, sizeof(int));
    }
    else
    {
        read(p[0], &y, sizeof(int));
        printf("%d %d", x, y);
    }
}
```

(5)

16 points

Problem II

- Using the SIGALRM signal, write a program that draws a number between 1 and 100 after one second and displays it.
- Write a C program under UNIX where the parent process creates 100 child processes as follows: The father creates the first process and waits for the child to display its PID, before creating the second child. Then he waits for the second to display his PID before creating the third child, and so on.


16 points

Problem III

- We consider a memory managed by contiguous allocation. The allocation of processes is made according to the first fit algorithm. That is, the first encountered area whose size is greater than or equal to the size of the process to be loaded is the one that is allocated

Page 1/2

the process. At the instant t , the state of the main memory is described in the figure below (where A, B, C, D, E, F are processes already loaded in memory, and the other parts are free zones):



Represent the evolution of the main memory according to the arrival of each of the following successive events.

- Event 1. Arrival of program G (20k).
- Event 2. Start of program B. \Rightarrow not eliminated that's not in memory
- Event 3. Arrival of program H (15k).
- Event 4. Departure of the program E. \Rightarrow not in memory
- Event 5. Arrival of program I (40k).

- The memory of a computer contains 4-page frames and, at the beginning, all the frames are empty. How many page faults produces the following page references 3, 4, 4, 1, 5, 2, 3, 1, 4 using, respectively, FIFO, OPTIMAL, and LRU replacement algorithms? Justify your answers by showing the contents of the frames after each reference.

20 points

Problem IV

- Describe (in pseudocode) the steps to be followed by the function `void Move (char * file_name, int f1, int f2)` which moves a file from a source directory (f1 is the index of its descriptor entry) to a destination directory (f2 is the index of its descriptor entry).
- Describe (in pseudocode) the steps to be followed by the function `int Similar (int f, char * file_name_1, file_name_2)` which compares the contents of two files belonging to the opened directory with entry descriptor f. If both files have the same content, the function returns 1, and 0 if not.

P.S: Feel free to use the structures and functions seen in the course.

Bon travail